

MSAGLJS: Browsing Large Graphs with Tile Pyramids and Sleeve Routing

Anonymous author(s)

Anonymous affiliation(s)

Abstract

We present a new way to visualize a large graph in the style of online geographic maps. The method builds a *tile pyramid* for *semantic zoom*: at every zoom level the labels of the highest-ranked nodes remain readable, just as the names of major geographical features stay readable on those maps.

The edges are routed by a method we call *sleeve routing*, which computes for each edge the shortest path in its homotopy class around the node obstacles, using the Constrained Delaunay Triangulation. We apply several heuristics to speed up the routing.

We implemented our approach in the WebGL renderer of MSAGLJS, an open-source TypeScript library for graph visualization in web browsers, with the entire pipeline running client-side, without a dedicated server. We target graphs with at most 40K nodes and 100K edges. The renderer's demo can be seen at <https://microsoft.github.io/msagljs/renderer-webgl/index.html>.

MSAGLJS is available on GitHub¹ and as NPM packages².

2012 ACM Subject Classification Human-centered computing → Graph drawings; Theory of computation → Computational geometry

Keywords and phrases Graph visualization, edge routing, Constrained Delaunay Triangulation, tiling, WebGL

Category Track 2, Long Paper

Generative AI Declaration Generative AI was used for editorial assistance in the preparation of this article.

1 Introduction

Visualizing graphs in web browsers without a dedicated server presents unique challenges: the runtime is single-threaded, the per-tab JavaScript heap is capped at roughly 4 GB,³ and users expect the smooth pan-and-zoom experience familiar from online maps. MSAGLJS⁴ is an open-source TypeScript library that addresses these challenges by combining efficient layout algorithms, a novel edge routing method, and a tiling scheme with semantic zoom.

MSAGLJS is consumed as a set of NPM packages and renders graphs using WebGL through the `deck.gl` framework [2]. It supports the layered Sugiyama scheme [44], Pivot MDS [14], and IPSep-CoLa [21], with edges routed around node obstacles. Throughout this paper our typical example is a social network laid out by IPSep-CoLa; the methods we describe, however, are agnostic to the layout algorithm and apply to any node placement in which the nodes do not overlap. The rendering pipeline builds a hierarchy of tiles—analogueous to map tiles—so that at any zoom level only a bounded number of entities are drawn. Tiles are delivered to the browser through the standard `TileLayer` of the `@deck.gl/geo-layers`

¹ <https://github.com/microsoft/msagljs>

² `@msagl/core`, `@msagl/drawing`, `@msagl/parser`, `@msagl/renderer-svg`, `@msagl/renderer-webgl` — all on <https://www.npmjs.com/>.

³ In current Chrome, V8 with pointer compression caps the per-tab JavaScript heap at approximately 4 GB (2^{32} bytes).

⁴ <https://github.com/microsoft/msagljs>



© Anonymous author(s);

licensed under Creative Commons License CC-BY 4.0

34th International Symposium on Graph Drawing and Network Visualization (GD 2026).

Editors: TBD; Article No. 0; pp. 0:1–0:12



Leibniz International Proceedings in Informatics

LIPIC Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

package, giving smooth pan and zoom. Live demos are available online: a WebGL renderer for large graphs⁵ and an SVG renderer for smaller, editable graphs⁶.

This paper makes two main contributions:

1. A *tiling scheme* for large graph visualization that builds a hierarchical tile pyramid, limits visible entities per tile using PageRank-based filtering, simplifies edge routes at coarser levels, and integrates with the WebGL renderer for real-time pan and zoom.
2. A *sleeve routing* algorithm that routes edges on the dual graph of a Constrained Delaunay Triangulation, using per-source batched Dijkstra trees followed by the classical funnel algorithm to produce shortest-path geodesics in homotopy classes.

2 Related Work

2.0.0.1 Graph visualization tools.


Graphviz [25, 5] is a long-standing graph drawing tool supporting multiple layout algorithms, including the Sugiyama method and Scalable Force-Directed Placement [7]; it does not support tiling or WebGL rendering. Our library builds on the earlier MSAGL desktop layout engine [39] that did not have the features described here. On the web, Sigma.js [8] and Cytoscape.js [23] provide interactive graph rendering, but rely on straight-line or generic spline edges without obstacle-avoiding routing and without a pyramid of precomputed tiles. ReGraph [6] uses WebGL to render large graphs with straight-line edges and supports interactive cluster expansion but not tiling. Cosmograph [1] uses GPU-accelerated force-directed layout for million-node graphs with straight-line edges, without tiling. GraphMaps [38] supports tiling and level-of-detail for large graphs but runs only on Windows and does not optimize edge routes. Cornac [41] handles huge graphs with tiles and edge bundling [34, 32] but requires a multi-server backend. Earlier navigation schemes for large graphs include ASK-GraphView [12], which organizes nodes into a hierarchical cluster tree, and topological fisheye views [24], which distort a multilevel layout [27, 33] under the focus; neither targets browser-side tile pyramids or per-level edge rerouting. Hierarchical edge bundling [31] and force-directed bundling [32] are complementary clutter-reduction techniques applied *after* edges have been routed.

2.0.0.2 Shortest paths and edge routing.

Our sleeve algorithm walks the dual of a Constrained Delaunay Triangulation [43] and relies on the classical funnel algorithm for extracting a geodesic from a triangle strip [36, 15, 26, 4]. The optimal algorithm for Euclidean shortest paths among polygonal obstacles [30] and the recent $O(m \log^{2/3} n)$ single-source shortest-path algorithm that breaks the sorting barrier [20] are remarkable theoretical results, but they are not practical: both build elaborate global data structures that we avoid. Graphviz builds the full visibility graph for edge routing in force-directed layouts, which has $O(n^2)$ edges. For Sugiyama layouts, it uses the funnel algorithm [18] but requires a simple containing polygon. yWorks [11] offers “Organic edge routing” based on force-directed simulation. The approach of Dwyer and Nachmanson [22] routes on a Yao-graph spanner with local shortcutting. For orthogonal layouts, the libavoid framework of Wybrow, Marriott and Stuckey performs obstacle-avoiding connector routing with incremental updates [45]. Our sleeve method eliminates the spanner entirely and routes

⁵ <https://microsoft.github.io/msagljs/renderer-webgl/index.html>

⁶ <https://microsoft.github.io/msagljs/renderer-svg/index.html>



124 ■ **Figure 1** The entities rendered on each level as the user zooms in shown without scaling.

73 directly on the CDT dual graph, and is, to our knowledge, the first routing algorithm
 74 integrated with a browser-side tile pyramid that reroutes edges per level.

75 2.0.0.3 CDT-based pathfinding.

76 The closest precedent to our routing pipeline is the CDT-and-funnel pathfinder of Demyen
 77 and Buro [17], originally designed for game-AI units of nonzero radius. They run admissible
 78 A* search on the CDT dual (with a reduced-graph variant that collapses corridor chains into
 79 single abstract edges) to recover provably shortest paths. Their search is, however, per-query:
 80 the bounds depend on the start and goal, so it cannot share work across the edges with the
 81 same source as our per-source Dijkstra does. For the graphs we target this would be roughly
 82 an order of magnitude slower; we therefore trade their proven optimality for speed.

83 3 Tiling

84 The input to tiling is a laid-out graph together with routed edges: we prefer to route the
 85 edges with sleeve routing (Section 4), to achieve the visual stability on a level change. The
 86 output is a sequence of $Z+1$ levels: level Z is the finest, and level 0 is the coarsest. Each level
 87 is a collection of tiles indexed by integer (x, y) grid coordinates. The output is consumed by
 88 the `TileLayer` of the `@deck.gl/geo-layers` package [10].

89 A tile is a rectangular viewport together with its *tile data*: nodes, edge labels, edge
 90 arrowheads, and *edge clips*. An edge clip is a curve confined to the tile rectangle, paired
 91 with the non-empty list of graph edges whose routes coincide with that curve inside the tile.
 92 Bundles capture the stretches on which several routes travel tight together inside a common
 93 sleeve; representing each such stretch as a single clip with an attached edge list, rather than
 94 as one clip per edge, lowers each tile's element count and the number of draw calls at render
 95 time, and avoids the extra subdivisions that an inflated count would otherwise trigger.

123 Pseudocode 1 describes the algorithm building the tile pyramid.

125 The single tile at level 0 has a rectangle equal to the graph bounding box. Each coarser

96 **Algorithm 1** BUILD_TILE_PYRAMID($G, \mathcal{C}, \text{minTileSize}, \mathcal{M}_{\max}$)

```

97 1:  $T_0 \leftarrow$  single tile holding all of  $G$  clipped to its bounding box
98 2:  $levels \leftarrow [T_0]$ ;  $z \leftarrow 0$ ;  $used \leftarrow |T_0|$   $\triangleright |T|$  = element count of tile  $T$ 
99 3: while some tile in  $levels[z]$  exceeds  $\mathcal{C}$  elements do
100 4:    $next \leftarrow$  empty  $2^{z+1} \times 2^{z+1}$  tile grid
101 5:   for all tiles  $T \in levels[z]$  do  $\triangleright$  per-tile inner loop
102 6:     split  $T$  into its four sub-tiles of  $next$   $\triangleright$  Section 3.2
103 7:      $used +=$  number of new elements added to  $next$ 
104 8:     if  $200 \cdot used > \mathcal{M}_{\max}$  then  $\triangleright$  memory budget; checked mid-split
105 9:       discard  $next$  and break pyramid growth
106 10:    end if
107 11:  end for
108 12:  if  $next$ 's tile width or height is below  $\text{minTileSize}$  then
109 13:    break  $\triangleright$  tile size below threshold
110 14:  end if
111 15:   $levels[z+1] \leftarrow next$ ;  $z \leftarrow z + 1$ 
112 16: end while
113 17:  $Z \leftarrow z$ ;  $V_Z \leftarrow V(G)$ ;  $s_v^{(Z)} \leftarrow 1$  for all  $v \in V_Z$ 
114 18: compute PageRank( $G$ )
115 19: for  $z = Z - 1$  down to 0 do
116 20:    $(V_z, \{s_v^{(z)}\}) \leftarrow \text{SELECT\_TOP\_K\_WITH\_ADAPTIVE\_SCALE}(V_{z+1}, 2^{Z-z})$   $\triangleright$  Section 3.1
117 21:   build a CDT from the inflated polylines of  $V_z$  at scales  $\{s_v^{(z)}\}$ 
118 22:   route every edge  $e \in E_z$  on this CDT, optionally using the level- $(z+1)$  route as a
119   sleeve hint  $\triangleright$  Section 4.2
120 23:   rebuild the tiles of  $levels[z]$   $\triangleright$  Section 3.2
121 24: end for
122 25: return  $levels$ 

```

126 level keeps a prefix of the PageRank-ordered nodes together with all edges whose endpoints
127 both lie in that prefix; the finest level keeps the full node and edge sets.

128 3.0.0.1 How the index of the finest layer, Z , is chosen.

129 We treat the tiles differently when calculating Z and when preparing for rendering. There is
130 no element filtering in the former case: all graph entities are represented in each level. We
131 continue increasing Z , generating new levels, and split tiles, until a stopping condition is
132 met. Then, while keeping the finest level intact, we rebuild the tiles of levels $0, \dots, Z - 1$ for
133 rendering.

134 We start from $Z = 0$ and grow the pyramid one level at a time. We stop as soon as any
135 of three conditions is met:

- 136 1. every tile at the new level contains at most $\mathcal{C}$ elements, with default $\mathcal{C} = 500$;
- 137 2. the new tile width and height are both below ten times the average node width and
138 height;
- 139 3. while generating the tiles of the new finest level, the running total of stored tile elements
140 summed over all levels, multiplied by 200 bytes per element, exceeds the memory budget
141 $\mathcal{M}_{\max}$, with default 4 GB; the partial level is then discarded and Z is set to the previous
142 finest level.

143 The third condition is checked incrementally inside the level-build loop, so we abort as




Figure 2 Two adjacent pyramid levels are shown, cropped to the same window; the coarser level is on the left, the finer one on the right. For each level a new graph is assembled from a PageRank-ranked subset of the nodes.

soon as the memory barrier is reached rather than finishing a level we know we will throw away. We additionally cap $Z \leq 8$ to bound the depth of the pyramid; the nine-level limit is comfortably above what is reachable on any input we tested.

3.0.0.2 Three phases.

Once Z and the tile rectangles are fixed, each coarser level $z = Z-1, Z-2, \dots, 0$ is built in three phases:

1. pick the nodes of level z with PageRank and possibly enlarge them, Section 3.1;
2. reroute the edges of G_z around the resulting obstacles, Section 4;
3. recompute the per-tile edge clips on levels $0, \dots, z$ by re-seeding the root tile with the new curves and applying the top-down midline split, Section 3.2.

Figure 2 shows two adjacent pyramid levels built by this procedure; Algorithm 1 summarizes the whole pipeline.

3.1 PageRank-Guided Level Construction

Before filling the level's tiles for rendering we run PageRank [40] on G once, and sort all nodes of the graph, V , by descending rank. Write $k = Z - z$ for the level's depth below the finest level. The candidates for level of depth k are the nodes of the *prefix* of the sorted array of length $\lceil |V|/2^k \rceil$.

Node positions are inherited from the original layout and never move: across levels we change only each node's display size.

We add nodes to a level of depth k by walking the level's prefix in rank-descending order. The top-ranked node is scaled by 2^k . We keep the node scales non-increasing, but as large as possible, while making sure that each next node does not overlap the previously selected nodes. If even at its original size the candidate's bounding box would overlap an already-accepted node's scaled box, the candidate is dropped.

Dropping a relatively high-ranked candidate at a coarse level is less disruptive than it sounds. Consider Arya's node in the *Game of Thrones* graph: as Figure 1 shows, it is

dropped from the top layer because a higher-ranked node, Tyrion, has covered the available region. The user still sees that part of the graph as occupied, and Arya’s individual node appears at a finer level where its box fits.

3.1.0.1 Spatial-hash overlap test.

A naive overlap test against every previously accepted node would dominate the cost of the filter. We instead maintain the accepted boxes in a uniform spatial hash whose cell size is chosen so that any accepted box, and any candidate’s box at the level’s maximum scale, both fit inside one cell on each axis. Because of this choice, two such boxes can overlap only if their center cells differ by at most one step on each axis: testing a candidate therefore reduces to scanning the accepted boxes registered in the candidate’s own cell and in the eight cells immediately around it.

The hash brings little benefit at the coarsest levels, where the prefix is short and even a naive scan over the few accepted boxes is cheap; it pays off at finer levels, where the prefix grows geometrically and a per-candidate constant-time test replaces a linear scan over hundreds of already-accepted boxes. On the layouts we tested the expected work per candidate is constant, so the overall cost is dominated by the initial sort by PageRank. In our experiments this phase was never a bottleneck: the per-level rerouting dominates the build time at every level we measured.

3.2 Splitting Edges between Tiles

We speed up the edge clip generation by implementing the following idea. Consider a tile and an edge clip with the curve that intersects the boundary of the tile only at curve’s endpoints. Cut the curve with the middle lines of the tile and consider the new clips obtained this way. They have the same property: all intersections of the smaller tile boundaries and the new edge clips happen at the endpoints.

When we create the level’s pyramid, line 6 of Algorithm 1, we have each edge clip already contained in its tile, so the subdivision starts from it. However, when we prepare the tiles for rendering, line 23 of Algorithm 1, we have graph G_z with rerouted edges of the level, so for each edge of G_z we start the subdivision by putting the edge curve in an empty tile containing the whole graph.

3.2.0.1 Bundling edge clips.

To correctly reflect the number of elements in a tile and speed up the edge clip calculation even more, clips within a tile whose two endpoints coincide are bundled into a single clip whose attached edge list accumulates every contributing edge; the bundled geometry is taken from the first occurrence. Substituting one short curve for another with matching endpoints also adds visual stability. Table 1 reports the resulting maximum number of elements rendered in a single tile across the entire pyramid for a range of graphs.

4 Sleeve Routing on the CDT

To prepare for routing given the graph layout, each node is covered by a padded obstacle polygon. We compute a Constrained Delaunay Triangulation $\mathcal{T}$ on the set of these polygons [16, 19]. The *dual graph* D of $\mathcal{T}$ has one node per triangle and one edge for each pair of triangles sharing a side. The weight of an edge is the Euclidean distance between the centroids of its triangles.

Table 1 Maximum number of elements rendered in a single tile across the whole pyramid for several graphs (capacity $\mathcal{C} = 500$, $Z_{\max} = 8$). The capacity $\mathcal{C}$ only guides the choice of Z ; we have no tight analytical bound on the actual rendered per-tile element count, which is therefore reported empirically. “Max tile’s level” is the level on which the maximum tile was found. In the worst cases, **ca-HepPh** and **ca-CondMat**, the densest tile holds $6\text{--}8\times$ the nominal capacity $\mathcal{C}$.

Graph	$ V $	$ E $	Levels built	Max per tile	Max tile’s level
gameofthrones	407	2 639	5	324	4
composers	3 405	13 832	8	749	3
ca-GrQc	5 242	28 968	8	359	4
ca-HepTh	9 877	51 946	8	1 836	3
facebook_combined	4 039	88 234	9	928	5
ca-HepPh	12 008	236 978	9	3 643	3
ca-CondMat	23 133	186 878	8	2 949	3
deezer_europe	28 281	92 752	8	2 402	3
delaunay_n15	32 768	98 274	8	283	7

4.1 Routing inside a Sleeve

We split routing an edge into two subproblems: a *combinatorial* one—which sequence of diagonals to cross—and a *geometric* one—which exact polyline through that sequence is shortest. The geometric step is the classical funnel algorithm [15, 29], which pulls the path taut inside the polygon formed by the sleeve triangles.

To obtain the sleeve we run a shortest-path search on D , where triangles inside obstacles other than the source and target obstacles are excluded. The natural per-edge instantiation is A* [28] with the straight-line distance to t as heuristic. When a node has many incident edges, however, running independent searches is wasteful. We instead group edges by source s and run a single Dijkstra computation per source, expanding until all targets are reached and recovering each sleeve by following parent pointers. This gives a speedup compared with A* routing, for example on the Game of Thrones social network (407 nodes, 2 639 edges) it cuts routing time by 30%.

We can improve even more. Since the query edges are undirected on D , each can be served from either endpoint, so the number of distinct roots is in turn reducible by reorienting query edges to minimize that count—an instance of minimum vertex cover on the demand graph, well approximated in linear time by the greedy maximum-degree rule and used as a source-side speed-up in many-to-many shortest-path frameworks [35]. Applying this heuristic reduces the number of Dijkstra trees from 331 (one per distinct source) to 199 on Game of Thrones and from 2 645 to 1 285 on the composers graph (3 405 nodes, 13 832 edges), shortening the routing pass by $\sim 17\%$ and $\sim 35\%$ respectively.

In addition, before launching the sleeve search for an edge we try routing the straight line segment between the source and the target by walking the CDT.

4.2 Edge Hint

We explored a routing variant in which the path generated for an edge on one level provides a hint for routing the same edge on an adjacent level, increasing the visual stability of the routes between levels. Its shortcoming is that the search cannot be batched per source as the Dijkstra-tree variant of Section 4.1 is, which impacts performance.




Figure 3 Effect of vertex collapse on the LORAS-RENTLY edge of the Game of Thrones graph. Endpoints are shown as blue dots; their padded obstacles are highlighted in light blue. (a) The CDT triangles forming the sleeve are drawn in cornflower blue and the funnel path in red; the thick black polyline is the perimeter of the non-collapsed sleeve. The yellow dot marks the source-obstacle vertex that gets collapsed. (b) The sleeve after collapse, drawn as a single purple polygon with the source/target obstacles merged into the corresponding endpoints, and the funnel path (red) through it. Collapsing widens the funnel opening at the endpoint and removes the spurious detour around the obstacle corner.

4.3 Vertex Collapse at Source and Target

We improve the sleeve's geometry in the situation depicted in Figure 3. The remedy is to *collapse* corner-hugging chain vertices to the corresponding endpoint. Walking the right chain of the sleeve outward from s , we find the first vertex of the source obstacle at which the chain turns right toward s , and replace every source-owned vertex on that chain up to and including this one by s . The left chain is processed mirror-symmetrically (with left turns instead of right), and the target end likewise. Diagonals whose two endpoints coincide after collapse are dropped, and the left/right orientation of each surviving diagonal is preserved from the uncollapsed geometry. The collapsed vertex in Figure 3(a) is marked in yellow.

After collapse, the funnel sees a wide opening at each endpoint, as in Figure 3(b), and naturally selects the optimal exit direction. A final arrowhead-trimming step clips the path at the node boundary curves. Collapse might cause the resulting path to overlap nodes other than the source and the target, but we have found the effect to be negligible.

The implementation also offers an optional pass that fits a Bezier segment into each polyline corner to render edges as smooth curves. It is disabled by default for performance: the polyline rendering is fast and visually acceptable, so we trade the smoother curves for routing speed.

5 Rendering with deck.gl

The output of the build pipeline of Section 3 is consumed in the browser by the `TileLayer` of `@deck.gl/geo-layers` [10]. `TileLayer` is the standard machinery that powers slippy-map renderers: it watches the viewport, decides which tiles are needed at the current zoom, and asynchronously requests them through a user-supplied `getTileData` callback, caching results across frames so that subsequent pans within the cache stay free of work.

`MSAGLJS` plugs into this machinery at construction time by exposing the precomputed pyramid as the `TileLayer`'s zoom range, with the coarsest pyramid level mapped to the layer's lowest zoom and the finest level to the highest. As the user pans and zooms, `TileLayer` resolves the viewport to an integer zoom and a set of axis-aligned tile indices and, for each tile it wants to keep visible, calls `getTileData`. The renderer translates each such request back into a tile of the corresponding pyramid level and returns its precomputed nodes, edge clips, edge labels, and arrowheads. The mapping is one-to-one, so no geometric work happens at view time: panning slides the tile cache, and zooming swaps in tiles from the adjacent level under `TileLayer`'s built-in cross-fade.

Because the routes baked into the tiles were already rerouted per level by the sleeve router of Section 4, every level the viewer lands on is visually self-consistent: the displayed edges are the geodesics around exactly the obstacles drawn on that level, not a subsampling of finer routes.

5.0.0.1 Interaction.

The renderer supports pan, zoom, and hover highlighting of nodes and edges. Hovering an edge highlights it and its incident nodes; hovering a node highlights the node, its one-hop neighbors, and its two-hop neighbors, each in a different color. A `zoomTo` API animates the viewport to a target rectangle with a smooth transition and is used, for example, by the search box to focus the view on a query result.

6 Experimental Results

We evaluate the sleeve router and the tile-pyramid build on the same benchmark graphs as Table 1: the Game of Thrones social network [13], the GD 2011 contest `composers` graph [9], four SNAP collaboration networks (`ca-GrQc`, `ca-HepTh`, `ca-HepPh`, `ca-CondMat`) and the SNAP `facebook_combined` ego network [37, 3], the `deezer_europe` social network [42], and the SuiteSparse Delaunay mesh `delaunay_n15`. All experiments run in Node.js (single process, headless; not a browser) on an Apple M4 Max laptop with 48 GB of unified memory. Layout uses `IPsepCola`, and the tile pyramid is built with the same parameters as Table 1, capacity $C = 500$ and $Z_{\max} = 8$.

Table 2 reports per-graph timings of three pipeline stages: *Routing*, the sum of the two phases logged by the sleeve router on the full graph—constrained Delaunay triangulation of the obstacles and the Dijkstra-tree search on the CDT dual (Section 4.1); *Tiling*, the build of the tile pyramid (Section 3), which includes the per-level rerouting; and *Total*, the end-to-end loading time including parsing.

The CDT phase contributes negligibly to the routing column—from 8 ms on Game of Thrones to about half a second on the largest graphs—so virtually all of *Routing* is the Dijkstra-tree sleeve search, with the vertex-cover heuristic of Section 4.1 reducing the number of trees we run. Even on graphs with hundreds of thousands of edges (`ca-CondMat`, `ca-HepPh`, `deezer_europe`) the full pipeline finishes in a few minutes; this is a one-time build cost. The

Graph	$ V $	$ E $	Routing	Tiling	Total
gameofthrones	407	2 639	0.17	0.17	1.87
composers	3 405	13 832	4.46	2.43	7.67
ca-GrQc	5 242	28 968	4.96	2.72	9.16
ca-HepTh	9 877	51 946	27.60	9.62	40.51
ca-HepPh	12 008	236 978	83.40	35.61	128.10
facebook_combined	4 039	88 234	13.82	6.55	22.33
ca-CondMat	23 133	186 878	202.50	55.29	273.60
deezer_europe	28 281	92 752	388.60	68.90	479.10
delaunay_n15	32 768	98 274	31.23	12.94	68.40

Table 2 End-to-end loading benchmarks on an Apple M4 Max laptop, all times in seconds. *Routing* sums the CDT-construction and Dijkstra-tree sleeve-search phases of Section 4.1 on the full graph; *Tiling* is the build of the tile pyramid (capacity $C = 500$, $Z_{\max} = 8$) of Section 3, which includes the per-level rerouting; *Total* is the total loading time including parsing of the graph file, graph layout, and all other stages.

resulting tile pyramid already contains the per-level routes, so panning and zooming only fetch and render the precomputed tiles, with no rerouting at view time.

7 Conclusion

We presented MSAGLJS, an open-source TypeScript library for graph layout, edge routing, and visualization in web browsers. The sleeve routing method routes edges directly on the CDT dual graph using batched Dijkstra trees and the funnel algorithm. The tiling scheme enables map-like exploration of large graphs with level-of-detail filtering and edge simplification. The WebGL renderer, built on deck.gl, provides smooth pan-and-zoom interaction.

MSAGLJS is available at <https://github.com/microsoft/msagljs> under the MIT license.

References

- 1 Cosmograph. <https://cosmograph.app>.
- 2 deck.gl: Large-scale WebGL-powered data visualization. <https://deck.gl/>.
- 3 facebookcombined. https://snap.stanford.edu/data/facebook_combined.txt.gz.
- 4 Funnel algorithm. <https://page.mi.fu-berlin.de/mulzer/notes/alggeo/polySP.pdf>.
- 5 Graphviz. <http://www.graphviz.org/>.
- 6 Regraph. <https://cambridge-intelligence.com/regraph/>.
- 7 sfdp. <https://graphviz.org/docs/layouts/sfdp/>.
- 8 Sigma.js: a JavaScript library aimed at visualizing graphs of thousands of nodes and edges. <https://www.sigmajs.org/>.
- 9 Skewed. <http://mozart.diei.unipg.it/gdcontest/contest2011/composers.xml>.
- 10 TileLayer (@deck.gl/geo-layers). <https://deck.gl/docs/api-reference/geo-layers/tile-layer>. Accessed: 2026.
- 11 yworks. <https://yworks.com/products/yed>.
- 12 James Abello, Frank Van Ham, and Neeraj Krishnan. ASK-GraphView: A large scale graph visualization system. In *IEEE Transactions on Visualization and Computer Graphics*, volume 12, pages 669–676. IEEE, 2006.
- 13 Andrew Beveridge and Michael Chemers. The game of game of thrones: Networked concordances and fractal dramaturgy. In *Reading Contemporary Serial Television Universes*, pages 201–225. Routledge, 2018.

- 373 14 Ulrik Brandes and Christian Pich. Eigensolver methods for progressive multidimensional
374 scaling of large data. In *Graph Drawing: 14th International Symposium, GD 2006, Karlsruhe,*
375 *Germany, September 18-20, 2006. Revised Papers 14*, pages 42–53. Springer, 2007.
- 376 15 Bernard Chazelle. A theorem on polygon cutting with applications. In *23rd Annual Symposium*
377 *on Foundations of Computer Science (sfcs 1982)*, pages 339–349. IEEE, 1982.
- 378 16 Boris Delaunay et al. Sur la sphere vide. *Izv. Akad. Nauk SSSR, Otdelenie Matematicheskii i*
379 *Estestvennyka Nauk*, 7(1):793–800, 1934.
- 380 17 Douglas Demyen and Michael Buro. Efficient triangulation-based pathfinding. In *Proceedings*
381 *of the 21st National Conference on Artificial Intelligence (AAAI)*, pages 942–947, 2006.
- 382 18 David P Dobkin, Emden R Gansner, Eleftherios Koutsofios, and Stephen C North. Implement-
383 ing a general-purpose edge router. In *Graph Drawing: 5th International Symposium, GD'97*
384 *Rome, Italy, September 18–20, 1997 Proceedings 5*, pages 262–271. Springer, 1997.
- 385 19 Vid Domiter and Borut Žalik. Sweep-line algorithm for constrained delaunay triangulation.
386 *International Journal of Geographical Information Science*, 22(4):449–462, 2008.
- 387 20 Ran Duan, Jiayi Mao, Xiao Mao, Xinkai Shu, and Longhui Yin. Breaking the sorting barrier
388 for directed single-source shortest paths. In *Proceedings of the 57th Annual ACM Symposium*
389 *on Theory of Computing (STOC)*, 2025. <https://arxiv.org/abs/2504.17033>.
- 390 21 Tim Dwyer, Yehuda Koren, and Kim Marriott. IPSep-CoLa: An incremental procedure for
391 separation constraint layout of graphs. *IEEE Transactions on Visualization and Computer*
392 *Graphics*, 12(5):821–828, 2006. doi:10.1109/TVCG.2006.156.
- 393 22 Tim Dwyer and Lev Nachmanson. Fast edge-routing for large graphs. In *Graph Drawing:*
394 *17th International Symposium, GD 2009, Chicago, IL, USA, September 22-25, 2009. Revised*
395 *Papers 17*, pages 147–158. Springer, 2010.
- 396 23 Max Franz, Christian T. Lopes, Gerardo Huck, Yue Dong, Onur Sumer, and Gary D. Bader.
397 Cytoscape.js: a graph theory library for visualisation and analysis. *Bioinformatics*, 32(2):309–
398 311, 2016.
- 399 24 Emden R. Gansner, Yehuda Koren, and Stephen North. Topological fisheye views for visualizing
400 large graphs. *IEEE Transactions on Visualization and Computer Graphics*, 11(4):457–468,
401 2005.
- 402 25 Emden R. Gansner and Stephen C. North. An open graph visualization system and its
403 applications to software engineering. *Software: Practice and Experience*, 30(11):1203–1233,
404 2000.
- 405 26 Leonidas Guibas, John Hershberger, Daniel Leven, Micha Sharir, and Robert E. Tarjan.
406 Linear-time algorithms for visibility and shortest path problems inside triangulated simple
407 polygons. *Algorithmica*, 2(1–4):209–233, 1987.
- 408 27 David Harel and Yehuda Koren. A fast multi-scale method for drawing large graphs. In
409 *Journal of Graph Algorithms and Applications*, volume 6, pages 179–202, 2002.
- 410 28 Peter E. Hart, Nils J. Nilsson, and Bertram Raphael. A formal basis for the heuristic
411 determination of minimum cost paths. *IEEE Transactions on Systems Science and Cybernetics*,
412 4(2):100–107, 1968. doi:10.1109/TSSC.1968.300136.
- 413 29 John Hershberger and Jack Snoeyink. Computing minimum length paths of a given homotopy
414 class. *Computational geometry*, 4(2):63–97, 1994.
- 415 30 John Hershberger and Subhash Suri. An optimal algorithm for Euclidean shortest paths in
416 the plane. *SIAM Journal on Computing*, 28(6):2215–2256, 1999.
- 417 31 Danny Holten. Hierarchical edge bundles: Visualization of adjacency relations in hierarchical
418 data. *IEEE Transactions on Visualization and Computer Graphics*, 12(5):741–748, 2006.
- 419 32 Danny Holten and Jarke J. Van Wijk. Force-directed edge bundling for graph visualization.
420 In *Computer Graphics Forum*, volume 28, pages 983–990. Wiley Online Library, 2009.
- 421 33 Yifan Hu. Efficient, high-quality force-directed graph drawing. *Mathematica Journal*, 10(1):37–
422 71, 2005.

- 423 **34** Christophe Hurter, Ozan Ersoy, and Alexandru Telea. Graph bundling by kernel density
424 estimation. In *Computer graphics forum*, volume 31, pages 865–874. Wiley Online Library,
425 2012.
- 426 **35** Sebastian Knopp, Peter Sanders, Dominik Schultes, Falk Schieferdecker, and Dorothea Wagner.
427 Computing many-to-many shortest paths using highway hierarchies. In *Proceedings of the 9th*
428 *Workshop on Algorithm Engineering and Experiments (ALENEX)*, pages 36–45. SIAM, 2007.
- 429 **36** D. T. Lee and Franco P. Preparata. Euclidean shortest paths in the presence of rectilinear
430 barriers. *Networks*, 14(3):393–410, 1984.
- 431 **37** Jure Leskovec, Jon Kleinberg, and Christos Faloutsos. Graph evolution: Densification and
432 shrinking diameters. *ACM transactions on Knowledge Discovery from Data (TKDD)*, 1(1):2–es,
433 2007.
- 434 **38** Lev Nachmanson, Roman Prutkin, Bongshin Lee, Nathalie Henry Riche, Alexander E Holroyd,
435 and Xiaoji Chen. Graphmaps: Browsing large graphs as interactive maps. In *Graph Drawing*
436 *and Network Visualization: 23rd International Symposium, GD 2015, Los Angeles, CA, USA,*
437 *September 24-26, 2015, Revised Selected Papers 23*, pages 3–15. Springer, 2015.
- 438 **39** Lev Nachmanson, George G. Robertson, and Bongshin Lee. Drawing graphs with GLEE. In
439 *Graph Drawing: 15th International Symposium, GD 2007*, pages 389–394. Springer, 2008.
- 440 **40** Lawrence Page, Sergey Brin, Rajeev Motwani, and Terry Winograd. The pagerank citation
441 ranking: Bringing order to the web. *Stanford InfoLab*, 249(373):1–17, 1999.
- 442 **41** Alexandre Perrot and David Auber. Cornac: Tackling huge graph visualization with big data
443 infrastructure. *IEEE Transactions on Big Data*, 6(1):80–92, 2018.
- 444 **42** Benedek Rozemberczki and Rik Sarkar. Characteristic Functions on Graphs: Birds of a
445 Feather, from Statistical Descriptors to Parametric Models. In *Proceedings of the 29th ACM*
446 *International Conference on Information and Knowledge Management (CIKM '20)*, page
447 1325–1334. ACM, 2020.
- 448 **43** Jonathan Richard Shewchuk. Delaunay refinement algorithms for triangular mesh generation.
449 *Computational Geometry*, 22(1–3):21–74, 2002.
- 450 **44** Kozo Sugiyama, Shojiro Tagawa, and Mitsuhiro Toda. Methods for visual understanding
451 of hierarchical system structures. *IEEE Transactions on Systems, Man, and Cybernetics*,
452 11(2):109–125, 1981. doi:10.1109/TSMC.1981.4308636.
- 453 **45** Michael Wybrow, Kim Marriott, and Peter J. Stuckey. Orthogonal connector routing. *Lecture*
454 *Notes in Computer Science*, 5849:219–231, 2010.