

Browsing large graphs with MSAGLJS, a graph dragh drawing tool in JavaScript

Lev Nachmanson and Xiaoji Chen

Microsoft Research, US,
levnach@hotmail.com, cxiaoji@gmail.com,
Msagljs github home page: <https://github.com/microsoft/msagljs>

Abstract. There has been progress in visualization of large graphs recently. Tools appeared that can render a huge graph in seconds. However, if we request that the node labels were rendered, and the edges were not overlapping the nodes they are not adjacent to, then the problem is still standing. Interacting with a large graph in an Internet browser with the same ease as browsing an online map, inspecting the high level structure and zooming in to the high level detail, is still a challenging task. In this paper we describe novel approaches to several aspects of this problem. We give a new algorithm for edge routing, where the edges do not overlap the nodes to which they are not adjacent. The algorithm produces edge paths which are visually appealing and optimal in their homotopy class. To facilitate graph visualization with DeckGL, we propose a new simple and fast tiling method. The method guarantees that in every view, except of the highest layer, the number of visible entities is not larger than a predefined bound. Our method provides a high level overview of the graph, with the gradual increase of the detail level. We make the node labels of the most important nodes for the current view visible. The edge routing algorithm mentioned above is reused at the tiling stage to simplify the paths on the lower levels. In addition, we bundle edges per-tile as an optimization heuristic

19 Introduction

20 We target our approach to large but not huge graphs. The maximum number of
21 vertices of the graphs we looked at was 28k, and the maximum number of the
22 edges was 237k. There are quite a few algorithms that calculate node positions
23 for such graphs, and work very fast [1, 2]. We look at the node layout as a solved
24 problem.

25 In the first part of the paper we address edge routing where an edge does
26 not intersects the nodes it is not adjacent to. Our approach works for any node
27 layout, as long as it produces a layout without overlap. We build on the edge
28 routing from [3] and improve it. There has been progress in visualization of
29 large graphs recently. Tools appeared that can render a huge graph in seconds.

30 However, the situation changes if we request that the node labels are rendered,
31 and the edges overlap only the nodes they are adjacent to. Interacting with a
32 large graph in an Internet browser with the same ease as browsing an online
33 map, inspecting the high level structure and zooming in to the high level detail,
34 is still an unsolved problem. In this paper we describe novel approaches to several
35 aspects of this problem.

36 We propose a novel and efficient algorithm for edge routing, where each edge
37 can only intersect its source or target. The algorithm produces edge paths which
38 are visually appealing and even optimal in their homotopy class.

39 To facilitate graph visualization with DeckGL, we propose a new simple and
40 fast tiling method. The method guarantees that in every view, except of the
41 views of the Shighest layer, the number of visible entities is not larger than a
42 predefined bound. The method can be used in other viewers that support tiling.

43 Our method provides a high level overview of the graph.

44 The edge routing algorithm mentioned above is reused at the tiling stage to
45 simplify the paths on the lower levels. In addition, we bundle edges per-tile as
46 an optimization heuristic.

47 Related work

48 A popular graph drawing tool Graphviz [4] applies Scalable Force-Directed Place-
49 ment [5] for large graphs, with no support for tiling. Its edge routing for this case
50 builds the whole visibility graph. This can be very slow because the visibility
51 graph can have $O(n^2)$ edges, where n is the number of the nodes in the graph.
52 Interestingly, the funnel algorithm [6, 7], the last step of our approach, is used
53 in Graphviz for the edge routing in the Sugiyama layout. We are not aware of
54 any tool that integrates Graphviz and uses tiling as well.

55 yWorks [8] has method "Organic edge routing" that produces edge routes
56 around the nodes. We could find only a very general description of the method:
57 "The algorithm is based on a force directed layout paradigm. Nodes act as re-
58 pulsive forces on edges in order to guarantee a certain minimal distance between
59 nodes and edges. Edges tend to contract themselves. Using simulated annealing,
60 this finally leads to edge layouts that are calculated for each edge separately".
61 It seems the algorithm runs in $O(n + m)\log(n + m)$ time, where n is the number
62 of the nodes and m is the number of the edges.

63 ReGraph [9] uses WebGL as the viewing platform. It can render a large graph
64 using straight lines for the edges. The tool does not support tiling, but instead
65 the user interactively opens the node that is a cluster of nodes.

66 "graph-tool.skewed" [10] does not implement its own layout algorithms or
67 edge routing algorithms, but instead provides a nice wrapper around the algo-
68 rithms from other layout tools.

69 Circos [11] visualizes large graphs in a circular layout. It does not support
70 tiles.

71 Cosmograph [12] uses a GPU to calculate the layout of a graph and can
 72 handle a graph with a million nodes. It renders edges as straight lines. It does
 73 not support tiling.

74 Edge routing in MSAGLJS

85 The edge routing starts, as in [3], by building a spanner graph, an approximation
 86 of the full visibility graph, and then finding shortest paths on the spanner. The
 87 spanner, see Fig. 2, is built on a variation of a Yao graph, which was introduced
 88 independently by Flinchbaugh and Jones [13] and Yao [14]. This kind of graph
 89 is defined by the set of cones with the apices at the vertices. The cones have the
 90 same angle, usually in the form of $\frac{2\pi}{n}$, where n is a natural number. The family
 91 of cones with the apex at a specific vertex partition the plane as illustrated in
 92 Fig. 1. For each cone at most one edge is created connecting the cone apex with
 93 a vertex inside of the cone, so the graph has $O(n)$ edges where n is the number
 94 of vertices.

95
 99 The approach of [3] applies local optimizations to shorten an edge path.
 100 Namely, it tries to shortcut one vertex at a time from the path, as illustrated in
 101 Fig 3. To smoothen a path, it fits Bezier segments into the polyline corners by
 102 using a binary search to find a larger fitting segment, see Fig 4. While analyzing
 103 performance of the edge routing in MSAGLJS, we noticed, that for a graph with
 104 more than 1k nodes these heuristics sometime create a performance bottleneck
 105 in spite of using R-Trees[15].

106 In addition, when the naive shortcutting of polyline corners fails, the resulting
 107 path might remain not visually appealing, as shown in Fig. 3.

108 We replace these heuristics with a more precise and efficient optimization
 109 described below.

110 Path optimization

111 Remember that a simple polygon is a polygon without holes.

112 An application of the 'path in a simple polygon' optimization to edge routing
 113 is not a new approach. The authors of Graphvis used it [16], but only for
 114 hierarchical layouts, where a simple polygon, \mathcal{P} , containing the path is available.
 115 They write: "If \mathcal{P} does not contain holes ... we can apply a standard "funnel"
 116 algorithm [6, 7] for finding Euclidean shortest paths in a simple polygon". In
 117 general case, for a non-layered layout, they build the visibility graph which is
 118 very expensive.

119 In our settings we are able to find the polygon \mathcal{P} even for any layout. We
 120 drop the requirement that \mathcal{P} is simple. Indeed, to run the "funnel" algorithm
 121 one only needs a sleeve: a sequence of triangles, where each triangle shares an
 122 edge with its successor, and leading from the start to the end of the path. Let us
 123 show how to build polygon \mathcal{P} , create a sleeve, and produce an optimized path.



Fig. 1. Yao graph

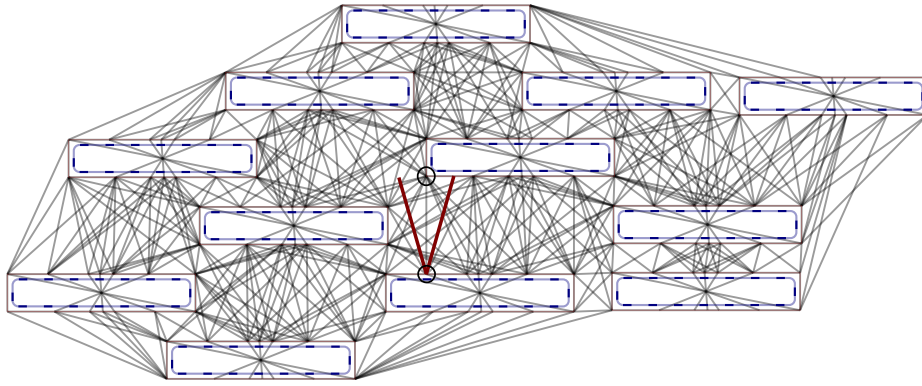


Fig. 2. Spanner graph is built using the idea of Yao graphs. The dashed curves are the original node boundaries. Each original curve is surrounded by a polygon with some offset to allow the polyline paths smoothing without intersecting the former. The edge marked by the circles is created because the top vertex is inside of the cone and it is the closest among such vertices to the cone apex. The apex of the cone is the lower vertex of the edge. MSAGLJS uses cone angle $\frac{\pi}{6}$, so the edges of the spanner can deviate from the optimal direction by this angle. Therefore, the shortest paths on the spanner have length that is at most the optimal shortest length multiplied by $\frac{1}{\cos(\frac{\pi}{6})} \simeq 1.155$.



Fig. 3. Unsuccessful shortcut


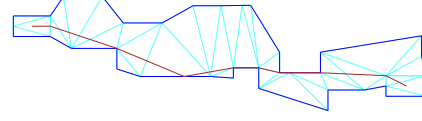


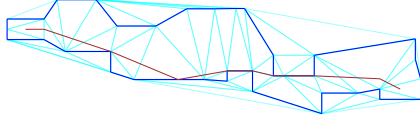
Fig. 4. Fitting a Bezier segment into a polyline corner




128 **Fig. 5.** Path \mathcal{L} with \mathcal{T} , a fragment.



129 **Fig. 6.** Polygon \mathcal{P} containing \mathcal{L} .



130 **Fig. 7.** New triangulation of \mathcal{P} .



131 **Fig. 8.** The optimized path together
132 with the sleeve diagonals.

124 We call obstacles, \mathcal{O} , the set of polygons covering the original nodes, see
125 Fig. 2. Before routing edges, we calculate a Constrained Delaunay Triangulation
126 [17] on \mathcal{O} . Let us call this triangulation \mathcal{T} .

127 For each edge of the graph we proceed with the following steps.

135 We route a path, called \mathcal{L} , on the spanner, as illustrated by Fig. 5. Let \mathcal{S} and
136 \mathcal{E} be the obstacles containing correspondingly \mathcal{L} 's start and end point. To obtain
137 \mathcal{P} , let us consider \mathcal{U} , the set of all triangles $t \in \mathcal{T}$ such that either $t \subset \mathcal{S} \cup \mathcal{E}$,
138 or t intersects \mathcal{L} and is not inside of any obstacle in $\mathcal{O} \setminus \{\mathcal{S}, \mathcal{E}\}$. The union of
139 \mathcal{U} gives us \mathcal{P} . The boundary of \mathcal{P} comprizes all edges e of the triangles from
140 \mathcal{U} such that e is adjacent to exactly one triangle from \mathcal{U} , see Fig. 6.

141 To create the sleeve [6, 7], we need to have a triangulation of \mathcal{P} such that every
142 edge of the triangulation is either a boundary edge of \mathcal{P} , or a diagonal of \mathcal{P} .
143 In our setup \mathcal{U} might not have this property, as in Fig. 6. We create a new
144 Constrained Delaunay Triangulation of \mathcal{P} , where the set of constrained edges is
145 the boundary of \mathcal{P} , see Fig. 7.

146 We trace path \mathcal{L} through the new triangulation and obtain the sleeve. Finally,
147 we apply the funnel algorithm on the sleeve and obtain the path which is the
148 shortest in the homotopy class of \mathcal{L} , as illustrated in Fig. 8.

149 The discussion [18] of the algorithm helped us in the implementation.

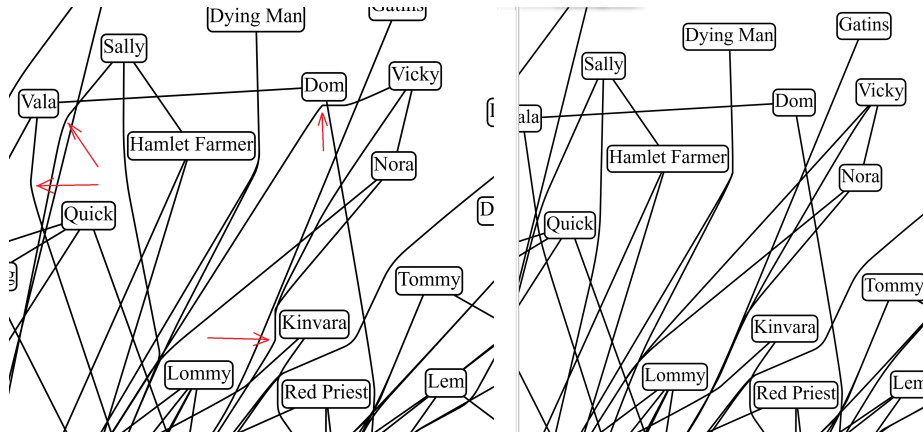
150 Polygon \mathcal{P} is not necessarily simple, as shown in Fig. 9. In this example the
151 path that we calculate with the funnel algorithm is not the shortest path inside
152 of \mathcal{P} .

153 Performance and quality comparison

157 In Fig. 10 we compare the paths generated by the old and the new method. We
158 can see that the paths produced by the new method have no kinks. We also
159 know that these paths are the shortest in their 'channels'. Arguably, the new
160 method produces better paths.



133 **Fig. 9.** \mathcal{P} is not simple. The dotted path is shorter than the dashed one that
 134 was found by the routing.



154 **Fig. 10.** The difference in the paths between the old, on the left, and the new,
 155 on the right, paths. The arrows on the left fragment point to the kinks that were
 156 removed by the new method.

Our performance experiments are summarized in Table. 1. We see that the older approach outperforms the new one on the smaller graphs; those with the number of nodes under 2000. The new method is faster on the rest of the graphs. We still prefer to use the new method independently of the graph size since the total slowdown is insignificant, under a half second in our experiments, but the quality of the paths is better. On the larger graphs the new method runs faster and produces better paths, so it is an obvious choice.

graph	nodes	edges	old method's time	new time
social network [19]	407	2639	1.0	1.4
b103 [20]	944	2438	1.6	2.0
b100 [21]	1463	5806	5.6	5.785
composers [22]	3405	13832	510.5	17.5
p2p-Gnutella04 [23]	10876	39994	375.4	293.8
facebook_combined [24]	4039	88234	132.2	119.1
lastfm_asia_edges [25]	7626	27807	43.3	41.4
deezer_europe_edges [25]	28283	92753	1596.9	1209.3
ca-HepPh [26]	12008	237010	521.2	495.0

Table 1. Performance comparison with time in seconds.

1 Tiling

The algorithm works in two phases. The first phase builds more and more detailed levels with smaller tiles until no more tile subdivision is required. Then second phase goes from the higher to lower levels and finalizes the levels.

A tile is a pair of a rectangle and data (*rect*, *tile_data*). Keys to the tile hierarchy are in the form (i, j, z) , where z is the level index and pair (i, j) indicates the rectangle inside of the level. The initial, the tile with the largest rectangle on level 0 is represented by the triplet $(0, 0, 0)$. For $z = 1$ there are four tiles $(0, 0, 1)$, $(0, 1, 1)$, $(1, 0, 1)$ and $(1, 1, 1)$. Each tile (i, j, z) can be subdivided into four tiles of the same size one level higher: $(2i, 2j, z + 1)$, $(2i, 2j + 1, z + 1)$, $(2i + 1, 2j, z + 1)$, and $(2i + 1, 2j + 1, z + 1)$.

Each z -level is represented by a map $L(z)$, so $L(z)(i, j)$ gives us a specific tile. During the first phase we can discover some empty tiles which correspond to $L(z)(i, j)$ being not defined.

The tiling works when the edge routing is done, so each edge e has an associated curve $c(e)$. During the subdivision process we create pairs *curve clips*, (e, p) , where p is $c(e)$ or a continuous trimmed piece of $c(e)$. By construction we will have the property that for each curve clip (e, p) the curve p belong to the corresponding tile rectangle and it might touch the boundary of rectangle only at the endpoints of p .

One of the parameters controlling the algorithm is the number for tile capacity, \mathcal{C} , setting the upper limit on how many elements can be visible in one tile. The elements could be a curve clip, an arrowhead, a node, or a label. In our setting \mathcal{C} is set by default to 10000.

The first phase starts with $L(0) = \{(0, 0) \rightarrow \text{tile data}\}$: the map consisting of only one lowest tile, and the elements of *tile data* are all curve clips $e, c(e)$, all graph nodes, all edge labels, and all edge arrowheads. If the total number of these elements is less than \mathcal{C} then the first phase stops; this is the usual case for a small graph.

If it is not the case then the first phase continues working. Let us suppose that the current level is z . We denote by $C(i, j)$ the number of elements in $L(z)(i, j)$, in other words, the number elements crossing tile (i, j, z) .

For the minimal size of the tile we take $(8 \times w, 8 \times h)$, where w is the average width and h is the average height of the nodes of the graph. The algorithm starts after the edge routing is done, so each edge has a curve, an optional label, and arrowheads associated with it. The algorithm keeps a map from tilesInitially, we create one top level tile and

References

1. Y. Hu and L. Shi, “Visualizing large graphs,” *Wiley Interdisciplinary Reviews: Computational Statistics*, vol. 7, no. 2, pp. 115–136, 2015.
2. U. Brandes and C. Pich, “Eigensolver methods for progressive multidimensional scaling of large data,” in *Graph Drawing: 14th International Symposium, GD 2006, Karlsruhe, Germany, September 18–20, 2006. Revised Papers 14*, pp. 42–53, Springer, 2007.
3. T. Dwyer and L. Nachmanson, “Fast edge-routing for large graphs,” in *Graph Drawing: 17th International Symposium, GD 2009, Chicago, IL, USA, September 22–25, 2009. Revised Papers 17*, pp. 147–158, Springer, 2010.
4. “Graphviz.” <http://www.graphviz.org/>.
5. “sfdp.” <https://graphviz.org/docs/layouts/sfdp/>.
6. B. Chazelle, “A theorem on polygon cutting with applications,” in *23rd Annual Symposium on Foundations of Computer Science (sfcs 1982)*, pp. 339–349, IEEE, 1982.
7. J. Hershberger and J. Snoeyink, “Computing minimum length paths of a given homotopy class,” *Computational geometry*, vol. 4, no. 2, pp. 63–97, 1994.
8. “yworks.” <https://yworks.com/products/yed>.
9. “Regraph.” <https://cambridge-intelligence.com/regraph/>.
10. “Skewed.” <https://graph-tool.skewed.de>.
11. “Circos.” <http://circos.ca/>.
12. “Cosmograph.” <https://cosmograph.app>.
13. B. Flinchbaugh and L. Jones, “Strong connectivity in directional nearest-neighbor graphs,” *SIAM Journal on Algebraic Discrete Methods*, vol. 2, no. 4, pp. 461–463, 1981.
14. A. C.-C. Yao, “On constructing minimum spanning trees in k-dimensional spaces and related problems,” *SIAM Journal on Computing*, vol. 11, no. 4, pp. 721–736, 1982.

- 244 15. A. Guttman, “R-trees: A dynamic index structure for spatial searching,” in *Pro-*
245 *ceedings of the 1984 ACM SIGMOD international conference on Management of*
246 *data*, pp. 47–57, 1984.
- 247 16. D. P. Dobkin, E. R. Gansner, E. Koutsofios, and S. C. North, “Implementing a
248 general-purpose edge router,” in *Graph Drawing: 5th International Symposium,*
249 *GD’97 Rome, Italy, September 18–20, 1997 Proceedings 5*, pp. 262–271, Springer,
250 1997.
- 251 17. B. Delaunay, “Sur la sphere vide, bull. acad. science ussr vii: Class,” *Sci. Mat. Nat*,
252 pp. 793–800, 1934.
- 253 18. “Funnel algorithm.” <https://page.mi.fu-berlin.de/mulzer/notes/alggeo/polySP.pdf>.
- 254 19. A. Beveridge and M. Chemers, “The game of game of thrones: Networked con-
255 cordances and fractal dramaturgy,” in *Reading Contemporary Serial Television*
256 *Universes*, pp. 201–225, Routledge, 2018.
- 257 20. “b103.” [https://github.com/microsoft/automatic-graph-](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot)
258 [layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot).
- 259 21. “b100.” [https://github.com/microsoft/automatic-graph-](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot)
260 [layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot).
- 261 22. “Skewed.” <http://mozart.diei.unipg.it/gdcontest/contest2011/composers.xml>.
- 262 23. “p2p-gnutella04.” <https://snap.stanford.edu/data/p2p-Gnutella04.html>.
- 263 24. “facebookcombined.” https://snap.stanford.edu/data/facebook_combined.txt.gz.
- 264 25. B. Rozemberczki and R. Sarkar, “Characteristic Functions on Graphs: Birds of a
265 Feather, from Statistical Descriptors to Parametric Models,” in *Proceedings of the*
266 *29th ACM International Conference on Information and Knowledge Management*
267 *(CIKM ’20)*, p. 1325–1334, ACM, 2020.
- 268 26. J. Leskovec, J. Kleinberg, and C. Faloutsos, “Graph evolution: Densification
269 and shrinking diameters,” *ACM transactions on Knowledge Discovery from Data*
270 *(TKDD)*, vol. 1, no. 1, pp. 2–es, 2007.