

Browsing large graphs with MSAGLJS, a graph dragh drawing tool in JavaScript

Lev Nachmanson and Xiaoji Chen

Microsoft Research, US,
levnach@hotmail.com, cxiaoji@gmail.com,
Msagljs github home page: <https://github.com/microsoft/msagljs>

1 **Abstract.** There has been progress in visualization of large graphs re-
2 cently. Tools appeared that can render a huge graph in seconds. However,
3 if we request that the node labels were rendered, and the edges were not
4 overlapping the nodes they are not adjacent to, then the problem is still
5 standing. Interacting with a large graph in an Internet browser with the
6 same ease as browsing an online map, inspecting the high level structure
7 and zooming in to the high level detail, is still an unsolved problem.
8 In this paper we describe novel approaches to several aspects of this
9 problem.
10 We give a new algorithm for edge routing, where the edges do not overlap
11 the nodes to which they are not adjacent. The algorithm produces edge
12 paths which are visually appealing and optimal in their homotopy class.
13 To facilitate graph visualization with DeckGL, we propose a new simple
14 and fast tiling method. The method guarantees that in every view, except
15 of the highest layer, the number of visible entities is not larger than a
16 predefined bound.
17 Our method provides a high level overview of the graph, with the grad-
18 ual increase of the detail level. We make the node labels of the most
19 important nodes for the current view visible.
 The edge routing algorithm mentioned above is reused at the tiling stage
 to simplify the paths on the lower levels. In addition, we bundle edges
 per-tile as an optimization heuristic

20 Introduction

21 We target our approach to large but not huge graphs. The maximum number of
22 vertices of the graphs we looked at was 28k, and the maximum number of the
23 edges was 237k. There are quite a few algorithms that calculate node positions
24 for such graphs, and work very fast [1, 2]. We look at the node layout as a solved
25 problem.

26 In the first part of the paper we address edge routing where an edge does
27 not intersects the nodes it is not adjacent to. Our approach works for any node
28 layout, as long as it produces a layout whithout overlap. We build on the edge
29 routing from [3] and improve it. There has been progress in visualization of

30 large graphs recently. Tools appeared that can render a huge graph in seconds.
 31 However, the situation changes if we request that the node labels are rendered,
 32 and the edges overlap only the nodes they are adjacent to. Interacting with a
 33 large graph in an Internet browser with the same ease as browsing an online
 34 map, inspecting the high level structure and zooming in to the high level detail,
 35 is still an unsolved problem. In this paper we describe novel approaches to several
 36 aspects of this problem.

37 We propose a novel and efficient algorithm for edge routing, where each edge
 38 can only intersect its source or target. The algorithm produces edge paths which
 39 are visually appealing and even optimal in their homotopy class.

40 To facilitate graph visualization with DeckGL, we propose a new simple
 41 and fast tiling method. The method guarantees that in every view, except of
 42 the highest layer, the number of visible entities is not larger than a predefined
 43 bound. The method can be used in other viewers that support tiling.

44 Our method provides a high level overview of the graph.

45 The edge routing algorithm mentioned above is reused at the tiling stage to
 46 simplify the paths on the lower levels. In addition, we bundle edges per-tile as
 47 an optimization heuristic.

48 Related work

49 A popular graph drawing tool Graphviz [4] applies Scalable Force-Directed Place-
 50 ment [5] for large graphs, with no support for tiling. Its edge routing for this case
 51 builds the whole visibility graph. This can be very slow because the visibility
 52 graph can have $O(n^2)$ edges, where n is the number of the nodes in the graph.
 53 Interestingly, the funnel algorithm [6, 7], the last step of our approach, is used
 54 in Graphviz for the edge routing in the Sugiyama layout. We are not aware of
 55 any tool that integrates Graphviz and uses tiling as well. Tool yWorks [?] has
 56 method "Organic edge routing" that produces edge routes around the nodes.
 57 The performance of the method is not clear. We could find only a very general
 58 description of the method: "The algorithm is based on a force directed layout
 59 paradigm. Nodes act as repulsive forces on edges in order to guarantee a certain
 60 minimal distance between nodes and edges. Edges tend to contract themselves.
 61 Using simulated annealing, this finally leads to edge layouts that are calculated
 62 for each edge separately". It seems the algorithm runs in $O(n + m)\log(n + m)$
 63 time, where n is the number of the nodes and m is the number of the edges.

64 Tool ReGraph [8] uses WebGL as the viewing platform. It can render a large
 65 graph using straight lines for the edges. The tool does not support tiling, but
 66 instead expects the user interactively open the node.

67 Tool "graph-tool.skewed" [9] does not implement its own layout algorithms
 68 or edge routing algorithms, but instead provides a nice wrapper around the
 69 algorithms from other layout tools.

70 [10]

71 [11]

72 machine learning approach [12]

73 [13]

74 [14]

75 Edge routing

86 The edge routing starts, as in [3], by building a spanner graph, an approximation
87 of the full visibility graph. The spanner, see Fig. 2, is built on a variation of a
88 Yao graph, which was introduced independently by Flinchbaugh and Jones [15]
89 and Yao [16]. This kind of graph is defined by the set of cones with the apices at
90 the vertices. The cones have the same angle, usually in the form of $\frac{2\pi}{n}$, where n
91 is a natural number, and. The family of cones with the apex at a specific vertex
92 partition the plane as illustrated in Fig. 1. For each cone at most one edge is
93 created connecting the cone apex with a vertex inside of the cone, so the graph
94 has $O(n)$ edges where n is the number of vertices.

95
99 The approach of [3] first builds a polyline path through the spanner, then
100 applies some local modifications to shorten and smoothen the path. It tries to
101 shortcut a vertex iteratively, as illustrated in Fig 3. To smoothen it fits Bezier seg-
102 ments into the polyline corners, using the binary search to find the larger fitting
103 segments, see Fig 4. While analyzing performance of edge routing in MSAGLJS,
104 we noticed that for a graph with more than 1000 nodes these heuristics some-
105 times create a performance bottleneck in spite of using R-Trees[17].
106 In addition, when the naive shortcutting of polyline corners fails, the resulting
107 path is not visually appealing, as shown in Fig. 3.

108 We replace these heuristics with a more precise optimization.

109 Path optimization

110 Remember that a simple polygon is a polygon without holes.

111 An application of the 'path in a simple polygon' optimization is not a new
112 approach. The authors of [18] used it, but only for hierarchical layouts, where a
113 simple polygon, \mathcal{P} , containing the path is available. They write: "If \mathcal{P} does not
114 contain holes ... we can apply a standard "funnel" algorithm [6, 7] for finding
115 Euclidean shortest paths in a simple polygon". In general case, for a non-layered
116 layout, they build the visibility graph which is very expensive.

117 Here we drop the requirement that \mathcal{P} is simple. Indeed, to run the "funnel"
118 algorithm one only needs a sleeve: a sequence of triangles adjoined on an edge
119 and leading from the start to the end of the path. We show how to build polygon
120 \mathcal{P} , create a sleeve, and produce an optimized path, for any layout. Let us describe
121 our method.

122 We call obstacles \mathcal{O} the set of polygons covering the original nodes, see Fig. 2.
123 Before routing edges we calculate a Constrained Delaunay Triangulation [19] on
124 \mathcal{O} and call it \mathcal{T} . Then for each edge of the graph we proceed with the following
125 steps.



Fig. 1. Yao graph



Fig. 2. Spanner graph is built using the idea of Yao graphs. The dashed curves are the original node boundaries. Each original curve is surrounded by a polygon with some offset to allow the polyline paths smoothing without intersecting the former. The edge marked by the circles is created because the top vertex is inside of the cone and it is the closest among such vertices to the cone apex. The apex of the cone is the lower vertex of the edge. MSAGLJS uses cone angle $\frac{\pi}{6}$, so the edges of the spanner can deviate from the optimal direction by this angle. Therefore, the shortest paths on the spanner have length that is at most the optimal shortest length multiplied by $\frac{1}{\cos(\frac{\pi}{6})} \simeq 1.155$.

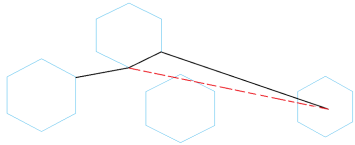


Fig. 3. Unsuccessful shortcut

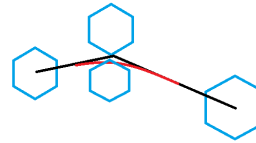
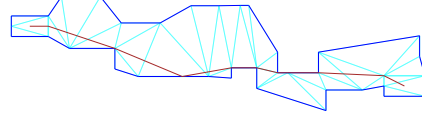


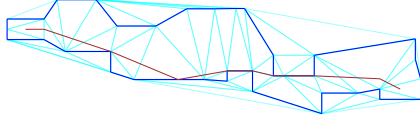
Fig. 4. Fitting a Bezier segment into a polyline corner




126 **Fig. 5.** Path \mathcal{L} with \mathcal{T} , a fragment.



127 **Fig. 6.** Polygon \mathcal{P} containing \mathcal{L} .



128 **Fig. 7.** New triangulation of \mathcal{P} .



129 **Fig. 8.** The optimized path together
130 with the sleeve diagonals.

133 We route a path, called \mathcal{L} , on the spanner, as illustrated by Fig. 5. Let \mathcal{S} and
134 \mathcal{E} be the obstacles containing correspondingly \mathcal{L} 's start and end point. To obtain
135 \mathcal{P} , let us consider \mathcal{U} , the set of all triangles $t \in \mathcal{T}$ such that either $t \subset \mathcal{S} \cup \mathcal{E}$,
136 or t intersects \mathcal{L} and is not inside of any obstacle in $O \setminus \{\mathcal{S}, \mathcal{E}\}$. The union of
137 \mathcal{U} gives us \mathcal{P} . The boundary of \mathcal{P} comprizes all edges e of the triangles from
138 \mathcal{U} such that e is adjacent to exactly one triangle from \mathcal{U} , see Fig. 6.

139 To create the sleeve [6, 7], we need to have a triangulation of \mathcal{P} such that every
140 edge of the triangulation is either a boundary edge of \mathcal{P} , or a diagonal of \mathcal{P} .
141 In our setup \mathcal{U} might not have this property, as in Fig. 6. We create a new
142 Constrained Delaunay Triangulation of \mathcal{P} , where the set of constrained edges is
143 the boundary of \mathcal{P} , see Fig. 7.

144 We trace \mathcal{L} through the new triangulation and obtain the sleeve. Finally, we
145 apply the funnel algorithm on the sleeve and obtain the path which is the shortest
146 in the homotopy class of \mathcal{L} , as illustrated in Fig. 8.

147 The discussion [20] of the algorithm helped us in the implementation.

148 Polygon \mathcal{P} is not necessarily simple, as shown in Fig. 9. In this example the
149 path that we calculate with the funnel algorithm is not the shortest path inside
150 of \mathcal{P} .

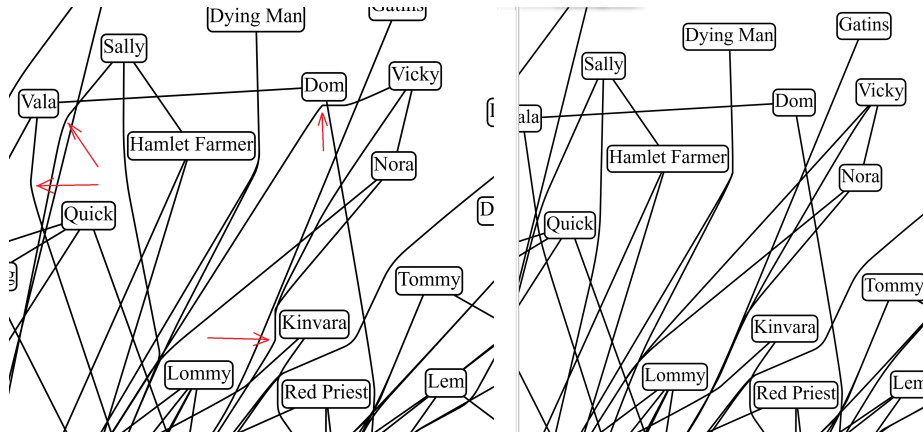
151 Performance and quality comparison

155 In Fig. 10 we compare the paths generated by the old and the new method. We
156 can see that the paths produced by the new method have no kinks. We also
157 know that these paths are the shortest in their 'channels'. Arguably, the new
158 method produces better paths.

170 Our performance experiments are summarized in Table. 1. We see that the
171 older approach outperforms the new one on the smaller graphs; those with the
172 number of nodes under 2000. The new method is faster on the rest of the graphs.
173 We still prefer to use the new method independently of the graph size since the
174 total slowdown is insignificant, under a half second in our experiments, but the



131 **Fig. 9.** \mathcal{P} is not simple. The dotted path is shorter than the dashed one that
 132 was found by the routing.



152 **Fig. 10.** The difference in the paths between the old, on the left, and the new,
 153 on the right, paths. The arrows on the left fragment point to the kinks that were
 154 removed by the new method.

quality of the paths is better. On the larger graphs the new method runs faster and produces better paths, so it is an obvious choice.

graph	nodes	edges	old method's time	new time
social network [21]	407	2639	1.0	1.4
b103 [22]	944	2438	1.6	2.0
b100 [23]	1463	5806	5.6	5.785
composers [24]	3405	13832	510.5	17.5
p2p-Gnutella04 [25]	10876	39994	375.4	293.8
facebook_combined [26]	4039	88234	132.2	119.1
lastfm_asia_edges [27]	7626	27807	43.3	41.4
deezer_europe_edges [27]	28283	92753	1596.9	1209.3
ca-HepPh [28]	12008	237010	521.2	495.0

Table 1. Performance comparison with time in seconds.

1 Tiling

The algorithm works in two phases. The first phase builds more and more detailed levels with smaller tiles until no more tile subdivision is required. Then second phase goes from the higher to lower levels and finalizes the levels.

A tile is a pair of a rectangle and data (*rect*, *tile_data*). Keys to the tile hierarchy are in the form (i, j, z) , where z is the level index and pair (i, j) indicates the rectangle inside of the level. The initial, the tile with the largest rectangle on level 0 is represented by the triplet $(0, 0, 0)$. For $z = 1$ there are four tiles $(0, 0, 1)$, $(0, 1, 1)$, $(1, 0, 1)$ and $(1, 1, 1)$. Each tile (i, j, z) can be subdivided into four tiles of the same size one level higher: $(2i, 2j, z + 1)$, $(2i, 2j + 1, z + 1)$, $(2i + 1, 2j, z + 1)$, and $(2i + 1, 2j + 1, z + 1)$.

Each z -level is represented by a map $L(z)$, so $L(z)(i, j)$ gives us a specific tile. During the first phase we can discover some empty tiles which correspond to $L(z)(i, j)$ being not defined.

The tiling works when the edge routing is done, so each edge e has an associated curve $c(e)$. During the subdivision process we create pairs *curve clips*, (e, p) , where p is $c(e)$ or a continuous trimmed piece of $c(e)$. By construction we will have the property that for each curve clip (e, p) the curve p belong to the corresponding tile rectangle and it might touch the boundary of rectangle only at the endpoints of p .

One of the parameters controlling the algorithm is the number for tile capacity, \mathcal{C} , setting the upper limit on how many elements can be visible in one tile. The elements could be a curve clip, an arrowhead, a node, or a label. In our setting \mathcal{C} is set by default to 10000.

201 The first phase starts with $L(0) = \{(0,0) \rightarrow \text{tile data}\}$: the map consisting
 202 of only one lowest tile, and the elements of *tile data* are all curve clips $e, c(e)$,
 203 all graph nodes, all edge labels, and all edge arrowheads. If the total number of
 204 these elements is less than \mathcal{C} then the first phase stops; this is the usual case for
 205 a small graph.

206 If it is not the case then the first phase continues working. Let us suppose
 207 that the current level is z . We denote by $C(i, j)$ the number of elements in
 208 $L(z)(i, j)$, in other words, the number elements crossing tile (i, j, z) .

209 For the minimal size of the tile we take $(8 \times w, 8 \times h)$, where w is the average
 210 width and h is the average height of the nodes of the graph. The algorithm starts
 211 after the edge routing is done, so each edge has a curve, an optional label, and
 212 arrowheads associated with it. The algorithm keeps a map from tilesInitially, we
 213 create one top level tile and

214 References

- 215 1. Y. Hu and L. Shi, “Visualizing large graphs,” *Wiley Interdisciplinary Reviews: Computational Statistics*, vol. 7, no. 2, pp. 115–136, 2015.
- 216 2. U. Brandes and C. Pich, “Eigensolver methods for progressive multidimensional scaling of large data,” in *Graph Drawing: 14th International Symposium, GD 2006, Karlsruhe, Germany, September 18-20, 2006. Revised Papers 14*, pp. 42–53, Springer, 2007.
- 217 3. T. Dwyer and L. Nachmanson, “Fast edge-routing for large graphs,” in *Graph Drawing: 17th International Symposium, GD 2009, Chicago, IL, USA, September 22-25, 2009. Revised Papers 17*, pp. 147–158, Springer, 2010.
- 218 4. “Graphviz.” <http://www.graphviz.org/>.
- 219 5. “sfdp.” <https://graphviz.org/docs/layouts/sfdp/>.
- 220 6. B. Chazelle, “A theorem on polygon cutting with applications,” in *23rd Annual Symposium on Foundations of Computer Science (sfcs 1982)*, pp. 339–349, IEEE, 1982.
- 221 7. J. Hershberger and J. Snoeyink, “Computing minimum length paths of a given homotopy class,” *Computational geometry*, vol. 4, no. 2, pp. 63–97, 1994.
- 222 8. “Regraph.” <https://cambridge-intelligence.com/regraph/>.
- 223 9. “Skewed.” <https://graph-tool.skewed.de>.
- 224 10. “Circos.” <http://circos.ca/>.
- 225 11. H. Gibson, J. Faith, and P. Vickers, “A survey of two-dimensional graph layout techniques for information visualisation,” *Information visualization*, vol. 12, no. 3-4, pp. 324–357, 2013.
- 226 12. O.-H. Kwon, T. Crnovrsanin, and K.-L. Ma, “What would a graph look like in this layout? a machine learning approach to large graph visualization,” *IEEE transactions on visualization and computer graphics*, vol. 24, no. 1, pp. 478–488, 2017.
- 227 13. Z. Lin, N. Cao, H. Tong, F. Wang, U. Kang, and D. H. Chau, “Interactive multi-resolution exploration of million node graphs,” in *IEEE VIS*, 2013.
- 228 14. “Cosmograph.” <https://cosmograph.app>.
- 229 15. B. Flinchbaugh and L. Jones, “Strong connectivity in directional nearest-neighbor graphs,” *SIAM Journal on Algebraic Discrete Methods*, vol. 2, no. 4, pp. 461–463, 1981.

- 246 16. A. C.-C. Yao, "On constructing minimum spanning trees in k-dimensional spaces
247 and related problems," *SIAM Journal on Computing*, vol. 11, no. 4, pp. 721–736,
248 1982.
- 249 17. A. Guttman, "R-trees: A dynamic index structure for spatial searching," in *Pro-
250 ceedings of the 1984 ACM SIGMOD international conference on Management of
251 data*, pp. 47–57, 1984.
- 252 18. D. P. Dobkin, E. R. Gansner, E. Koutsofios, and S. C. North, "Implementing a
253 general-purpose edge router," in *Graph Drawing: 5th International Symposium,
254 GD'97 Rome, Italy, September 18–20, 1997 Proceedings 5*, pp. 262–271, Springer,
255 1997.
- 256 19. B. Delaunay, "Sur la sphere vide, bull. acad. science ussr vii: Class," *Sci. Mat. Nat*,
257 pp. 793–800, 1934.
- 258 20. "Funnel algorithm." <https://page.mi.fu-berlin.de/mulzer/notes/alggeo/polySP.pdf>.
- 259 21. A. Beveridge and M. Chemers, "The game of game of thrones: Networked con-
260 cordances and fractal dramaturgy," in *Reading Contemporary Serial Television
261 Universes*, pp. 201–225, Routledge, 2018.
- 262 22. "b103." [https://github.com/microsoft/automatic-graph-
263 layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot).
- 264 23. "b100." [https://github.com/microsoft/automatic-graph-
265 layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot).
- 266 24. "Skewed." <http://mozart.diei.unipg.it/gdcontest/contest2011/composers.xml>.
- 267 25. "p2p-gnutella04." <https://snap.stanford.edu/data/p2p-Gnutella04.html>.
- 268 26. "facebookcombined." https://snap.stanford.edu/data/facebook_combined.txt.gz.
- 269 27. B. Rozemberczki and R. Sarkar, "Characteristic Functions on Graphs: Birds of a
270 Feather, from Statistical Descriptors to Parametric Models," in *Proceedings of the
271 29th ACM International Conference on Information and Knowledge Management
272 (CIKM '20)*, p. 1325–1334, ACM, 2020.
- 273 28. J. Leskovec, J. Kleinberg, and C. Faloutsos, "Graph evolution: Densification
274 and shrinking diameters," *ACM transactions on Knowledge Discovery from Data
275 (TKDD)*, vol. 1, no. 1, pp. 2–es, 2007.