

Browsing large graphs with XJS, a graph drawing tool in JavaScript

auth0 and auth1

Institute, US,
 auth0@hotmail.com, auth1@gmail.com,
 X github home page: <https://github.com/X>

1 **Abstract.** There has been progress in visualization of large graphs re-
 2 cently. Tools appeared that can render a huge graph in seconds. However,
 3 if we request that the node labels are visible, and the edges are routed
 4 around the nodes, then the problem remains difficult. Interacting with
 5 a large graph in an Internet browser with the same ease as browsing
 6 an online map is still a challenging task. In this paper we describe a
 7 few novel approaches to large graph visualization that we developed in
 8 open-source JavaScript software.
 9 We give a new efficient edge routing algorithm, where the edges are
 10 routed around the nodes. The algorithm produces edge paths which are
 11 visually appealing and shortest in their homotopy class.
 12 To facilitate graph visualization with WebGL, or any other platform
 13 supporting tiles, we propose a new simple and efficient tiling method.
 14 The method guarantees that in every view, except of the highest level,
 15 the number of visible entities per tile is not larger than a predefined
 16 bound.
 17 We make the node labels of the most important nodes of the current
 18 view visible.
 The edge routing algorithm mentioned above is reused at the tiling stage
 to simplify the paths on the lower levels.

19 Introduction

20 Our software is open source, it is represented by a set of NPM packages. It
 21 runs on the client desktop or on a phone, and renders the graph in an Internet
 22 browser. We target large but not huge graphs. The maximum number of vertices
 23 of the graphs we applied our tool at was 28k, and the maximum number of the
 24 edges was 237k.

25 The algorithms described below were discovered while we programmed our
 26 tool. We believe these algorithms can be useful to other developers as well. The
 27 findings seemed to us interesting enough to put them into a paper.

28 The paper has sections Introduction, Related Work, Edge routing in XJS,
 29 Tiling, and Future work.

30 Let us start with a short review of some relevant to us publications.

31 Related work

32 A popular graph drawing tool Graphviz [1] applies method Scalable Force-
33 Directed Placement [2] for large graphs, with no support for tiling. Its edge
34 routing for this method builds the whole visibility graph and routes edges on it.
35 This can be very slow because the visibility graph can have $O(n^2)$ edges, where
36 n is the number of the nodes in the graph. Interestingly, the funnel algorithm [3,
37 4], the last step of our approach, is used in Graphviz for the edge routing in the
38 Sugiyama layout. We are not aware of any tool that integrates Graphviz and
39 uses tiling as well.

40 yWorks [5] has method "Organic edge routing" that produces edge routes
41 around the nodes. We could find only a very general description of the method:
42 "The algorithm is based on a force directed layout paradigm. Nodes act as re-
43 pulsive forces on edges in order to guarantee a certain minimal distance between
44 nodes and edges. Edges tend to contract themselves. Using simulated annealing,
45 this finally leads to edge layouts that are calculated for each edge separately".
46 It seems the algorithm runs in $O(n + m)\log(n + m)$ time, where n is the number
47 of the nodes and m is the number of the edges.

48 ReGraph [6] uses WebGL as the viewing platform. It can render a large graph
49 using straight lines for the edges. The tool does not support tiling, but instead
50 the user interactively opens the node that is a cluster of nodes.

51 "graph-tool.skewed" [7] does not implement its own layout algorithms or edge
52 routing algorithms, but instead provides a nice wrapper around the algorithms
53 from other layout tools.

54 Circos [8] visualizes large graphs in a circular layout. It does not support
55 tiles.

56 Cosmograph [9] uses a GPU to calculate the layout of a graph and can
57 handle a graph with a million nodes. It renders edges as straight lines. It does
58 not support tiling.


59 The authors of [10] implemented GraphMaps, a tool for large graph visu-
60 alization. The tool only runs on Windows. The edge were routed as polylines
61 on a triangulation and were not optimized. The tool supported tiling, but the
62 problem of the limiting number of visible entities was not solved.

63 In [11] an approach to visualize a huge graph is described. The method uses
64 tiles and edge bundling following [12], which is applied at the last moment during
65 the the graph browsing. The latter calculation is done on the client side. The
66 rest and the majority of the calculations runs on several servers.


67 Edge routing in XJS

78 The edge routing starts, as in [13], by building a spanner graph, an approximation
79 of the full visibility graph, and then finding shortest paths on the spanner. The
80 spanner, see Fig 2, is built on a variation of a Yao graph, which was introduced
81 independently by Flinchbaugh and Jones [14] and Yao [15]. This kind of graph
82 is defined by the set of cones with the apices at the vertices. The cones have the

83 same angle, usually in the form of $\frac{2\pi}{n}$, where n is a natural number. The family
 84 of cones with the apex at a specific vertex partition the plane as illustrated in
 85 Fig. 1. For each cone at most one edge is created connecting the cone apex with
 86 a vertex inside of the cone, so the graph has $O(n)$ edges where n is the number
 87 of vertices.



68 **Fig. 1.** Yao graph



69 **Fig. 2.** Spanner graph is built using the idea of Yao graphs. The dashed curves are the
 70 original node boundaries. Each original curve is surrounded by a polygon with some
 71 offset to allow the polyline paths smoothing without intersecting the former.
 72 The edge marked by the circles is created because the top vertex is inside the cone,
 73 and it is the closest among such vertices to the cone apex. The apex of the cone is the
 74 lower vertex of the edge.
 75 XJS uses cone angle $\frac{\pi}{6}$, so the edges of the spanner can deviate from the optimal
 76 direction by this angle. Therefore, the shortest paths on the spanner have length that
 77 is at most the optimal shortest length multiplied by $\frac{1}{\cos(\frac{\pi}{6})} \simeq 1.155$.

88
 92 The approach of [13] applies local optimizations to shorten an edge path.
 93 Namely, it tries to shortcut one vertex at a time from the path, as illustrated in
 94 Fig 3. To smoothen a path, it fits Bezier segments into the polyline corners by




Fig. 3. Unsuccessful shortcut




Fig. 4. Fitting a Bezier segment into a polyline corner

using a binary search to find a larger fitting segment, see Fig 4. While analyzing performance of the edge routing in XJS, we noticed, that for a graph with more than 1k nodes these heuristics sometime create a performance bottleneck, in spite of using R-Trees[16].

In addition, when the naive shortcutting of polyline corners fails, the resulting path might remain not visually appealing, as shown in Fig. 3.

We replace these heuristics with a more precise and efficient optimization described below.

Path optimization

We finalize edge routes by the “funnel” algorithm [3, 4], routing a path inside a simple polygon, that is a polygon without holes.


An application of the ‘path in a simple polygon’ optimization to edge routing is not a new idea: the novelty of our work is in how we find the polygon and how we use it. The authors of Graphvis used the ‘funnel’ algorithm [17], but only for hierarchical layouts, where a simple polygon, \mathcal{P} , containing the path is available. They write: “If \mathcal{P} does not contain holes ... we can apply a standard “funnel” algorithm ... for finding Euclidean shortest paths in a simple polygon”. In general case, for a non-layered layout, they build the visibility graph which is very expensive for a large graph.

Here we find the polygon \mathcal{P} for any layout. We drop the requirement that \mathcal{P} is simple. Indeed, to run the “funnel” algorithm one only needs a “sleeve”: a sequence of triangles leading from the start to the end of the path, where each triangle shares a side with its successor. Let us show how to build polygon \mathcal{P} , create a sleeve, and produce an optimized path.


We call obstacles, \mathcal{O} , the set of polygons covering the original nodes, see Fig. 2. Before routing edges, we calculate a Constrained Delaunay Triangulation [18] on \mathcal{O} . Let us call this triangulation \mathcal{T} .

For each edge of the graph we proceed with the following steps.


We route a path, called \mathcal{L} , on the spanner, as illustrated by Fig. 5. Let \mathcal{S} and \mathcal{E} be the obstacles containing correspondingly \mathcal{L} ’s start and end point. To obtain \mathcal{P} , let us consider \mathcal{U} , the set of all triangles $t \in \mathcal{T}$ such that either $t \subset \mathcal{S} \cup \mathcal{E}$, or t intersects \mathcal{L} and is not inside of any obstacle in $\mathcal{O} \setminus \{\mathcal{S}, \mathcal{E}\}$. The union of \mathcal{U} gives




123 **Fig. 5.** Path \mathcal{L} with \mathcal{T} , a fragment.



124 **Fig. 6.** Polygon \mathcal{P} containing \mathcal{L} .



125 **Fig. 7.** New triangulation of \mathcal{P} .



126 **Fig. 8.** The optimized path together
127 with the sleeve diagonals.

134 us \mathcal{P} . The boundary of \mathcal{P} comprizes all sides e of the triangles from \mathcal{U} such that
135 e belongs to exactly one triangle from \mathcal{U} , see Fig. 6.

136 To create the sleeve [3, 4], we need to have a triangulation of \mathcal{P} such that every
137 edge of the triangulation is either a boundary edge of \mathcal{P} , or a diagonal of \mathcal{P} .
138 Because \mathcal{U} might not have this property, as in Fig. 6, we create a new Constrained
139 Delaunay Triangulation of \mathcal{P} , where the set of constrained edges is the boundary
140 of \mathcal{P} , see Fig. 7.

141 We trace path \mathcal{L} through the new triangulation and obtain the sleeve. Finally,
142 we apply the funnel algorithm on the sleeve and obtain the path which is the
143 shortest in the homotopy class of \mathcal{L} , as illustrated in Fig. 8.

144 The discussion [19] of the algorithm helped us in the implementation.

145 Polygon \mathcal{P} is not necessarily simple, as shown in Fig. 9. In this example the
146 path that we calculate with the funnel algorithm is not the shortest path inside
147 of \mathcal{P} .

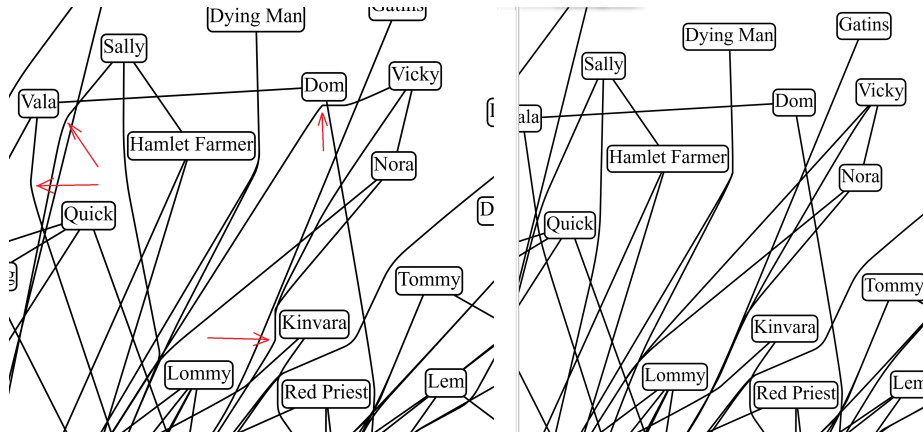
148 Performance and quality comparison

152 In Fig. 10 we compare the paths generated by the old and the new method. We
153 can see that the paths produced by the new method have no kinks. We also
154 know that these paths are the shorterst in their 'channels'. Arguably, the new
155 method produces better paths.

156 Our performance experiments are summarized in Table. 1. We see that the
157 older approach outperforms the new one on the smaller graphs; those with the
158 number of nodes under 2000. The new method is faster on the rest of the graphs.
159 We still prefer to use the new method independently of the graph size since
160 the slowdown is insignificant, but the quality of the paths is better. On the
161 larger graphs the new method runs faster and produces better paths, so it is an
162 obvious choice. To load a large graph, for example, `deezer_europe.edges` [20], we
163 start Edge or Chrome with an option that increases the memory limit of their
164 process: `- max_old_space_size=8192`.



128 **Fig. 9.** \mathcal{P} is not simple. The dotted path is shorter than the dashed one that
 129 was found by the routing.



149 **Fig. 10.** The difference in the paths between the old, on the left, and the new,
 150 on the right, paths. The arrows on the left fragment point to the kinks that were
 151 removed by the new method.

graph	nodes	edges	old method's time	new time
social network [21]	407	2639	1.0	1.4
b103 [22]	944	2438	1.6	2.0
b100 [23]	1463	5806	5.6	5.785
composers [24]	3405	13832	510.5	20.3
p2p-Gnutella04 [25]	10876	39994	375.4	304.2
facebook_combined [26]	4039	88234	132.2	123.7
lastfm_asia_edges [20]	7626	27807	43.3	54.7
deezer_europe_edges [20]	28283	92753	1596.9	1402.6
ca-HepPh [27]	12008	237010	521.2	495.0

Table 1. Performance comparison with time in seconds.

1 Tiling

The algorithm works in three phases. The first phase builds the levels starting from the lowest level and proceeding to higher and more detailed levels, with smaller tiles, until no more tile subdivision is required. The second phase filters out the entities from the layers to satisfy the capacity quota. Finally, the third phase simplifies the edge routes to utilize the space freed by the filtered out entities.

A tile, in our settings, is a pair $(rect, tiledata)$, where $rect$ is the rectangle of the tile and $tiledata$ is a set of *tile elements* visible in $rect$. A *tile element* could be a node, an edge label, an edge arrowhead, or an *edge clip*. An edge clip is a pair (e, p) , where e is an edge and p is a continuous piece of the edge curve c_e . Sometimes we need several edge clips to trace an edge through a tile.

The initial tile, the only tile on level 0, is represented by pair $(0, 0)$. For $z = 1$, there are four tiles: $(0, 0)$, $(0, 1)$, $(1, 0)$, and $(1, 1)$. Each tile (i, j) can be subdivided into four subtiles for level $z + 1$: $(2i, 2j)$, $(2i, 2j + 1)$, $(2i + 1, 2j)$, and $(2i + 1, 2j + 1)$.

Each z -level is represented by a map L_z , so $L_z(i, j)$ gives us a specific tile. Empty tiles correspond to undefined $L_z(i, j)$.


We use edge clips to represent the edge intersections with the tiles and provide the renderer with the minimal geometry that is sufficient to render a tile. To achieve this we require property \mathcal{F} :

a) For each tile t , for each curve clip $(e, p) \in t.tiledata$, we have: $p \subset t.rect$ and p might cross the boundary of the $t.rect$ only at endpoints of p .

b) For each edge e we have : the union of all p for all $(e, p) \in t.tiledata$ is equal to $c_e \cap t.rect$.

First phase of tiling

The first phase starts with $L_0 = \{(0, 0) \rightarrow (rect, tiledata)\}$: and $tiledata$ comprising curve clips (e, c_e) , for all edges e of the graph, all graph nodes, all edge



220 **Fig. 11.** Intersect curve $[start, end]$ with the midlines. Sort the intersections pa-
 221 rameters together with start, and end into array $u = [start, a, b, c, end]$. Split the
 222 curve to sub-curves $[start, a]$, $[a, b]$, $[b, c]$, $[c, end]$.

204 labels, and all edge arrowheads. We ensure property \mathcal{F} by setting $rect$ to a
 205 padded bounding box of the graph, so each edge curve does not intersect the
 206 boundary of $rect$.

207 Let us assume that L_z is already constructed and \mathcal{F} holds for its tiles. To
 208 build level L_{z+1} we divide each tile $t = L_z(i, j)$ into four subtiles of equal size.
 209 For each node, arrowhead, or edge label of $t.tiledata$, if the bounding box of the
 210 element intersects the subtile's rectangle then we add the element to the subtile
 211 $tiledata$.

212 The edge clip treatment is more involved. Let (e, p) be a curve clip belonging
 213 to tile t . We find all intersections of curve p with the horizontal midline and the
 214 vertical midline of $t.rect$. Each intersection can be represented as $p[t_j]$. We sort
 215 sequence $u = [start, \dots, t_j, \dots, end]$, where $[start, end]$ is the parameter domain
 216 of p , in ascending order, and remove the duplicates.

217 Next we create curve clips $(e, l_k) = (e, trim(p, u_k, u_{k+1}))$, as shown in Fig 11.
 218 We assign each curve clip (e, l_k) to the subtile with the rectangle containing the
 219 bounding box of l_k .

223 Because, by the induction assumption property \mathcal{F} is true on L_z , and by
 224 construction, each new curve clip can cross the boundary of the subtile only at
 225 the clip endpoints. We also cover all the intersections of p with the subtiles with
 226 the new edge clips, so the property \mathcal{F} holds for L_{z+1} .

227 Two parameters control the algorithm: tile capacity, \mathcal{C} , and the minimal size
 228 of a tile: $(\mathcal{W}, \mathcal{H})$. If for each (i, j) the number of elements in $L_z(i, j).tiledata$ is
 229 not greater than \mathcal{C} , or, if $w \leq \mathcal{W}$ and $h \leq \mathcal{H}$, where w (h) is the current tile
 230 width (correspondingly, height), then the second phase starts.

231 In our setting $\mathcal{C} = 500$, and $(\mathcal{W}, \mathcal{H}) = 3(w, h)$, where w is the average width
 232 and h is the average height of the nodes of the graph.

233 **Edge bundling** In our settings each edge clip is uniquely defined, module
 234 direction, by its start and end point. We can use this property to bundle the
 235 edges. In each tile we keep a map from unordered pairs of points to the set of
 236 edge clips that have these points as start and end points. Each such pair defines
 237 an edge bundle. For all edge clips in a bundle we create only one curve segment,
 238 avoiding the expensive trimming. We also count a bundle as one element in the
 239 tile, as in most of the cases the drawing attributes of the edges in the bundle are
 240 the same.

241 In our experiments, the number of edge bundles is about 50% of the number
 242 of edge clips, so the edge bundling is a significant optimization.

243 Second phase of tiling

244 The second phase of tiling filters out the entities from the lower layers. We do not
 245 change the highest, the most detailed layer. We sort the nodes of the graph into
 246 array N by PageRank [28]. For each layer L , except of the highest, we proceed
 as follows.

```

1: procedure FILTER( $L$ )
2:    $r \leftarrow \text{removeEntities}(L)$ 
3:   for all  $n$  in  $N$  do
4:     if ! $\text{addNodeToLayer}(n, r, N)$  then break
5:   end if
6: end for
7: end procedure

```

247 Here $\text{removeEntities}(L)$ empties all the tiles of layer L , but returns map r allow-
 248 ing to restore the tiles. Function $\text{addNodeToLayer}(n)$ returns false and does not
 249 change L when one of the tiles intersecting n already has more elements than
 250 \mathcal{C} . Otherwise, the function adds n to all tiles intersected by n . It also adds the
 251 tile elements for self edges of n , and the tile elements for the edges connecting n
 252 with the nodes appearing in N before n , i.e. the nodes with the rank not lesser
 253 than the rank of n .
 254

255 This procedure guarantees that each tile of L has no more than \mathcal{C} nodes, but
 256 a tile can have more than \mathcal{C} elements in general.

257 Third phase of tiling

258 In the third phase we use a fact that some nodes are not present on the layer. For
 259 all layers, except of the highest, we reroute the edges but only around the nodes
 260 that are present in the layer. We do not calculate edge routes from scratch, but
 261 use the existing routes and only apply the "funnel" heuristic in larger channels.
 262 This gives us simpler edge routes but still has a visual stability during the layer
 263 change while browsing.

264 2 Future work

- 265 – Find a tiling method that guarantees that each tile has no more than \mathcal{C} el-
266 ements. One approach could be to use a more aggressive edge bundling to
267 reduce the number of edge clips in the tiles.

268 References

- 269 1. “Graphviz.” <http://www.graphviz.org/>.
270 2. “sfdp.” <https://graphviz.org/docs/layouts/sfdp/>.
271 3. B. Chazelle, “A theorem on polygon cutting with applications,” in *23rd Annual*
272 *Symposium on Foundations of Computer Science (sfcs 1982)*, pp. 339–349, IEEE,
273 1982.
274 4. J. Hershberger and J. Snoeyink, “Computing minimum length paths of a given
275 homotopy class,” *Computational geometry*, vol. 4, no. 2, pp. 63–97, 1994.
276 5. “yworks.” <https://yworks.com/products/yed>.
277 6. “Regraph.” <https://cambridge-intelligence.com/regraph/>.
278 7. “Skewed.” <https://graph-tool.skewed.de>.
279 8. “Circos.” <http://circos.ca/>.
280 9. “Cosmograph.” <https://cosmograph.app>.
281 10. L. Nachmanson, R. Prutkin, B. Lee, N. H. Riche, A. E. Holroyd, and X. Chen,
282 “Graphmaps: Browsing large graphs as interactive maps,” in *Graph Drawing and*
283 *Network Visualization: 23rd International Symposium, GD 2015, Los Angeles, CA,*
284 *USA, September 24-26, 2015, Revised Selected Papers 23*, pp. 3–15, Springer, 2015.
285 11. A. Perrot and D. Auber, “Cornac: Tackling huge graph visualization with big data
286 infrastructure,” *IEEE Transactions on Big Data*, vol. 6, no. 1, pp. 80–92, 2018.
287 12. C. Hurter, O. Ersoy, and A. Telea, “Graph bundling by kernel density estimation,”
288 in *Computer graphics forum*, vol. 31, pp. 865–874, Wiley Online Library, 2012.
289 13. T. Dwyer and L. Nachmanson, “Fast edge-routing for large graphs,” in *Graph*
290 *Drawing: 17th International Symposium, GD 2009, Chicago, IL, USA, September*
291 *22-25, 2009. Revised Papers 17*, pp. 147–158, Springer, 2010.
292 14. B. Flinchbaugh and L. Jones, “Strong connectivity in directional nearest-neighbor
293 graphs,” *SIAM Journal on Algebraic Discrete Methods*, vol. 2, no. 4, pp. 461–463,
294 1981.
295 15. A. C.-C. Yao, “On constructing minimum spanning trees in k-dimensional spaces
296 and related problems,” *SIAM Journal on Computing*, vol. 11, no. 4, pp. 721–736,
297 1982.
298 16. A. Guttman, “R-trees: A dynamic index structure for spatial searching,” in *Pro-*
299 *ceedings of the 1984 ACM SIGMOD international conference on Management of*
300 *data*, pp. 47–57, 1984.
301 17. D. P. Dobkin, E. R. Gansner, E. Koutsofios, and S. C. North, “Implementing a
302 general-purpose edge router,” in *Graph Drawing: 5th International Symposium,*
303 *GD’97 Rome, Italy, September 18–20, 1997 Proceedings 5*, pp. 262–271, Springer,
304 1997.
305 18. B. Delaunay, “Sur la sphere vide, bull. acad. science ussr vii: Class,” *Sci. Mat. Nat*,
306 pp. 793–800, 1934.
307 19. “Funnel algorithm.” <https://page.mi.fu-berlin.de/mulzer/notes/alggeo/polySP.pdf>.

- 308 20. B. Rozemberczki and R. Sarkar, “Characteristic Functions on Graphs: Birds of a
309 Feather, from Statistical Descriptors to Parametric Models,” in *Proceedings of the*
310 *29th ACM International Conference on Information and Knowledge Management*
311 *(CIKM '20)*, p. 1325–1334, ACM, 2020.
- 312 21. A. Beveridge and M. Chemers, “The game of game of thrones: Networked con-
313 cordances and fractal dramaturgy,” in *Reading Contemporary Serial Television*
314 *Universes*, pp. 201–225, Routledge, 2018.
- 315 22. “b103.” [https://github.com/microsoft/automatic-graph-](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot)
316 [layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot).
- 317 23. “b100.” [https://github.com/microsoft/automatic-graph-](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot)
318 [layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot).
- 319 24. “Skewed.” <http://mozart.diei.unipg.it/gdcontest/contest2011/composers.xml>.
- 320 25. “p2p-gnutella04.” <https://snap.stanford.edu/data/p2p-Gnutella04.html>.
- 321 26. “facebookcombined.” https://snap.stanford.edu/data/facebook_combined.txt.gz.
- 322 27. J. Leskovec, J. Kleinberg, and C. Faloutsos, “Graph evolution: Densification
323 and shrinking diameters,” *ACM transactions on Knowledge Discovery from Data*
324 *(TKDD)*, vol. 1, no. 1, pp. 2–es, 2007.
- 325 28. L. Page, S. Brin, R. Motwani, and T. Winograd, “The pagerank citation ranking:
326 Bringing order to the web,” *Stanford InfoLab*, vol. 249, no. 373, pp. 1–17, 1999.