

New edge routing algorithm, and the tiling scheme in MSAGLJS

Lev Nachmanson and Xiaoji Chen

Microsoft Research, Redmond, US,
levnach@hotmail.com, cxiaoji@gmail.com,
MSAGLJS github page: <https://github.com/microsoft/MSAGLJS>

Abstract. One of the standard approaches to visualise a large amount of data with a computer is to present it with levels-of-details and use tiling, the way the online maps do. The tiles, often organized into a tree, are usually pre-calculated and dynamically loaded, according to the zoom level and the viewport. The root tile of the tree corresponds to the most simplified view, and the child tiles gradually give more details.

The first contribution of the paper is the tiling scheme of MSAGLJS. In this scheme we limit the number of entities visible per tile, except of the most detailed tiles.

To simplify the edge routes as we move from the most detailed to the less-detailed levels, we use a novel edge routing algorithm: our second, and main, contribution. The algorithm is efficient and produces visually appealing paths. More specifically, it routes the edges around the nodes and produces the shortest paths in their homotopy class. The latter property gives us a rather smooth change of the edge routes between the levels.

The third contribution is an efficient method finding the intersection of the edge path and the tile.

1 Introduction

Our open-source software, MSAGLJS, is available at GitHub at the address shown under the title. It is written in TypeScript and is usually consumed as a set of NPM packages. MSAGLJS facilitates rendering graphs on client computers, or phones, by working in an Internet browser. We target large but not huge graphs.

The paper consists of two main sections. The first section describes our edge routing algorithm that creates paths that are shortest in their homotopy class. We measure the performance of our new edge routing algorithm on public benchmark sets and compare it with the previous state-of-the-art routing algorithm.

The second section describes our new tiling method. The method allows loading graphs in a browser with up to 10K nodes and 40K edges and interacting with them in the style of online maps.

2 Related work

Popular graph drawing tool Graphviz [1] applies method Scalable Force-Directed Placement [2] for large graphs, with no support for tiling. The edge routing for this method builds the whole visibility graph and routes edges on it. This can be very slow because the visibility graph can have $O(n^2)$ edges, where n is the number of the nodes in the graph. Interestingly, the funnel algorithm [3, 4], the last step of our approach, is used in Graphviz for the edge routing in the Sugiyama layout. We are not aware of any tool that integrates Graphviz and uses tiling as well.

yWorks [5] has method "Organic edge routing" that produces edge routes around the nodes. We could find only a very general description of the method: "The algorithm is based on a force directed layout paradigm. Nodes act as repulsive forces on edges in order to guarantee a certain minimal distance between nodes and edges. Edges tend to contract themselves. Using simulated annealing, this finally leads to edge layouts that are calculated for each edge separately". It seems the algorithm runs in $O(n + m)\log(n + m)$ time, where n is the number of the nodes and m is the number of the edges.

ReGraph [6] uses WebGL as the viewing platform. It can render a large graph using straight lines for the edges. The tool does not support tiling, but instead the user interactively opens the node that is a cluster of nodes.

"graph-tool.skewed" [7] does not implement its own layout algorithms or edge routing algorithms, but instead provides a nice wrapper around the algorithms from other layout tools.

Circos [8] visualizes large graphs in a circular layout. It does not support tiles.

Cosmograph [9] uses a GPU to calculate the layout of a graph and can handle a graph with a million nodes. It renders edges as straight lines. It does not support tiling.

The authors of [10] implemented GraphMaps, a tool for large graph visualization. The tool only runs on Windows. The edge were routed as polylines on a triangulation and were not optimized. The tool supported tiling, but the problem of the limiting number of visible entities was not solved.

In [11] an approach to visualize a huge graph is described. The method uses tiles and edge bundling following [12], which is applied at the last moment during the graph browsing. The latter calculation is done on the client side. The rest and the majority of the calculations runs on several servers.

3 Edge routing

The user study of Xu et al [13] shows that straight edges improve the user comprehension of a graph drawing. From the other hand, Holten et al [14] show that strongly curved edges do not perform well in this regard. In our routing, we try to keep the edges as straight as possible, and to curve them just enough to

avoid the nodes. Our motivation for developing the routing described here was to improve the quality of edges and the algorithm speed.

The edge routing starts, as in [15], by building a spanner graph, an approximation of the full visibility graph, and then finding the edge routes as shortest paths on the spanner. The spanner, see Fig 2, is built on a variation of a Yao graph, which was introduced independently by Flinchbaugh, and Jones [16], and Yao [17]. This graph is built with a help of a set of cones with the apices at the vertices. Each cone of the set has the same angle, usually in the form of $\frac{2\pi}{k}$, where k is a natural number, $k = 12$ in our settings. The family of cones with the apex at a specific vertex partition the plane, as illustrated in Fig. 1. For each cone at most one edge is created connecting the cone apex with a vertex inside the cone. This way the spanner has at most kn edges, where n is the number of the vertices. We cover each node by a polygon with a relatively small number of corners, at most 8. Polygon corners play role of the vertices of the spanner. As a result, the spanner has $O(N)$ edges, where N is the number of the graph nodes.

The approach of [15] applies local optimizations to shorten an edge path. Namely, it tries to shortcut one vertex at a time from the path, as illustrated in Fig 3. To smoothen a path, it fits Bezier segments into the polyline corners by using a binary search to find a large fitting segment, see Fig 4. We noticed that when the shortcutting of polyline corners fails, the resulting path might remain not visually appealing, as shown in Fig. 3. We replace the shortcutting with a more precise, but still efficient optimization described below: that is one of the main contributions of our work.

3.1 Path optimization

We finalize edge routes by a slight modification of the “funnel” algorithm [3, 4], routing a path inside a simple polygon, that is a polygon without holes.

An application of the ‘path in a simple polygon’ optimization to edge routing is not a new idea: the novelty of our work is in how we find the polygon and how we use it. The authors of Graphviz used the ‘funnel’ algorithm [18], but only for hierarchical layouts, where a simple polygon, \mathcal{P} , containing the path is available. They write: “If \mathcal{P} does not contain holes ... we can apply a standard “funnel” algorithm ... for finding Euclidean shortest paths in a simple polygon”. In general case they build the visibility graph which is very expensive for a large graph.

Here we find the polygon \mathcal{P} for any layout. We drop the requirement that \mathcal{P} is simple. Indeed, to run the “funnel” algorithm one only needs a “sleeve”: a sequence of triangles leading from the start to the end of the path, where each triangle shares a side with its successor. Let us show how to build polygon \mathcal{P} , create a sleeve, and produce an optimized path.

We call obstacles, \mathcal{O} , the set of polygons covering the original nodes, see Fig. 2. Before routing edges, we calculate a Constrained Delaunay Triangulation [19] on \mathcal{O} , following [20]. Let us call this triangulation \mathcal{T} .

For each edge of the graph we proceed with the following steps.



Fig. 1: Yao graph




Fig. 2: Spanner graph is built using the idea of Yao graphs. The dashed curves are the original node boundaries. Each original curve is surrounded by a polygon with some offset to allow the polyline paths smoothing without intersecting the former.

The edge marked by the circles is created because the top vertex is inside the cone, and it is the closest among such vertices to the cone apex. The apex of the cone is the lower vertex of the edge.

MSAGLJS uses cone angle $\frac{\pi}{6}$, so the edges of the spanner can deviate from the optimal direction by this angle. Therefore, the shortest paths on the spanner have length that is at most the optimal shortest length multiplied by $\frac{1}{\cos(\frac{\pi}{6})} \simeq 1.155$.




Fig. 3: Unsuccessful shortcut




Fig. 4: Fitting a Bezier segment into a polyline corner




Fig. 5: Path \mathcal{L} with \mathcal{T} , a fragment.

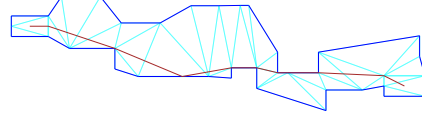


Fig. 6: Polygon \mathcal{P} containing \mathcal{L} .




Fig. 7: New triangulation of \mathcal{P} .




Fig. 8: The optimized path together with the sleeve diagonals.

We route a path, called \mathcal{L} , on the spanner, as illustrated by Fig. 5. Let \mathcal{S} and \mathcal{E} be the obstacles containing, correspondingly, \mathcal{L} 's start and end point. To obtain \mathcal{P} , let us consider \mathcal{U} , the set of all triangles $t \in \mathcal{T}$ such that either $t \subset \mathcal{S} \cup \mathcal{E}$, or t intersects \mathcal{L} and is not inside any obstacle in $\mathcal{O} \setminus \{\mathcal{S}, \mathcal{E}\}$. The union of \mathcal{U} gives us \mathcal{P} . The boundary of \mathcal{P} comprises all sides e of the triangles from \mathcal{U} such that e belongs to exactly one triangle from \mathcal{U} , see Fig. 6.

To create the sleeve [3, 4], we need to have a triangulation of \mathcal{P} such that every edge of the triangulation is either a boundary edge of \mathcal{P} , or a diagonal of \mathcal{P} . Because \mathcal{U} might not have this property, as in Fig. 6, we create a new Constrained Triangulation of \mathcal{P} , where the set of constrained edges is the boundary of \mathcal{P} , see Fig. 7.

We trace path \mathcal{L} through the new triangulation and obtain the sleeve. Finally, we apply the funnel algorithm on the sleeve and obtain the path which is the shortest in the homotopy class of \mathcal{L} , as illustrated in Fig. 8.

The discussion [21] of the algorithm helped us in the implementation of the funnel algorithm.

Polygon \mathcal{P} is not necessarily simple, as shown in Fig. 9. In this example the path that we calculate with the funnel algorithm is not the shortest path inside \mathcal{P} .

3.2 Performance and quality comparison

In Fig. 10 we compare the paths generated by the old and the new method. We can see that the paths produced by the new method have no kinks. We also know that these paths are the shortest in their 'channels'. Arguably, the new method produces better paths.

Our performance experiments are summarized in Table. 1. We see that the older approach outperforms the new one on the smaller graphs; those with the number of nodes under 2000. The new method is faster on the rest of the graphs.



Fig. 9: \mathcal{P} is not simple. The dotted path is shorter than the dashed one that was found by the routing.

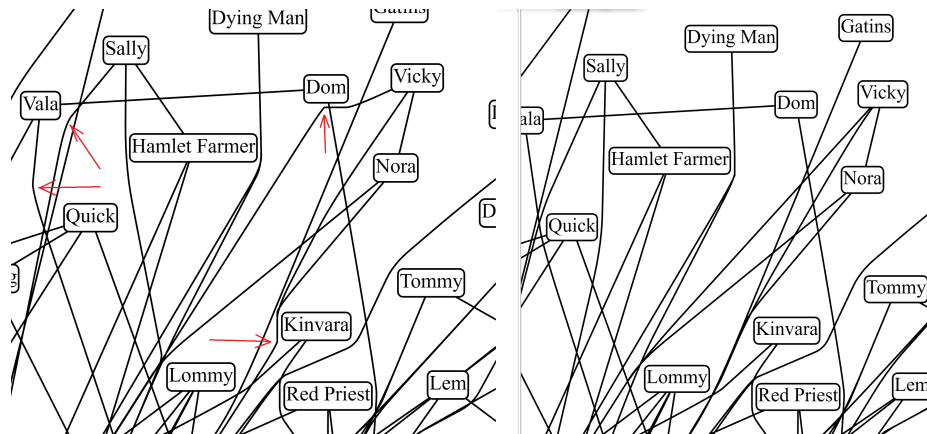


Fig. 10: Comparing the old, on the left, and the new, on the right, paths. The arrows on the left fragment point to the kinks that were removed by the new method.

We prefer the new method regardless of the graph size because it provides better path quality and the slowdown is insignificant.

graph	nodes	edges	old method's time	new time
social network [22]	407	2639	1.0	1.4
b103 [23]	944	2438	1.6	2.0
b100 [24]	1463	5806	5.6	5.785
composers [25]	3405	13832	510.5	20.3
p2p-Gnutella04 [26]	10876	39994	375.4	304.2
facebook_combined [27]	4039	88234	132.2	123.7
lastfm_asia_edges [28]	7626	27807	43.3	54.7
deezer_europe_edges [28]	28283	92753	1596.9	1402.6
ca-HepPh [29]	12008	237010	521.2	495.0

Table 1: Performance comparison with time in seconds.

4 Tiling

We had two goals when working on tiling. The first goal was to make exploring the graph in our tool similar to using online maps. The second goal was efficiency. The algorithm works in three phases. The first phase builds the levels starting from the lowest level and proceeding to higher and more detailed levels, with smaller tiles, until no more tile subdivision is required. The second phase filters out the entities from the layers to satisfy the capacity quota, as in [10]. Finally, the third phase simplifies the edge routes to utilize the space freed by the filtered out entities.

A tile, in our settings, is a pair $(rect, tiledata)$, where $rect$ is the rectangle of the tile and $tiledata$ is a set of *tile elements* visible in $rect$. A *tile element* could be a node, an edge label, an edge arrowhead, or an *edge clip*. An edge clip is a pair (e, p) , where e is an edge and p is a continuous piece of the edge curve c_e . Sometimes we need several edge clips to trace an edge through a tile.

The initial tile, the only tile on level 0, is represented by pair $(0, 0)$. For $z = 1$, there are four tiles: $(0, 0)$, $(0, 1)$, $(1, 0)$, and $(1, 1)$. Each tile (i, j) can be subdivided into four sub-tiles for level $z + 1$: $(2i, 2j)$, $(2i, 2j + 1)$, $(2i + 1, 2j)$, and $(2i + 1, 2j + 1)$.

Each z -level is represented by a map L_z , so $L_z(i, j)$ gives us a specific tile. Empty tiles correspond to undefined $L_z(i, j)$.

We use edge clips to represent the edge intersections with the tiles and provide the renderer with the minimal geometry that is sufficient to render a tile. To achieve this we require property \mathcal{F} :

a) For each tile t , for each edge clip $(e, p) \in t.tiledata$, we have: $p \subset t.rect$ and p might cross the boundary of the $t.rect$ only at endpoints of p .

b) For each edge e we have : the union of all p for all $(e, p) \in t.tiledata$ is equal to $c_e \cap t.rect$.

4.1 First phase of tiling

The first phase starts with $L_0 = \{(0, 0) \rightarrow (rect, tiledata)\}$: and *tiledata* comprising edge clips (e, c_e) , for all edges e of the graph, all graph nodes, all edge labels, and all edge arrowheads. We ensure property \mathcal{F} by setting *rect* to a padded bounding box of the graph, so each edge curve does not intersect the boundary of *rect*.

Let us assume that L_z is already constructed and \mathcal{F} holds for its tiles. To build level L_{z+1} we divide each tile $t = L_z(i, j)$ into four sub-tiles of equal size. For each node, arrowhead, or edge label of *t.tiledata*, if the bounding box of the element intersects the sub-tile's rectangle then we add the element to the sub-tile *tiledata*.

The edge clip treatment is more involved. Let (e, p) be an edge clip belonging to tile t . We find all intersections of curve p with the horizontal midline and the vertical midline of *t.rect*. Each intersection can be represented as $p[t_j]$. We sort sequence $u = [start, \dots, t_j, \dots, end]$, where $[start, end]$ is the parameter domain of p , in ascending order, and remove the duplicates.

Next we create edge clips $(e, l_k) = (e, trim(p, u_k, u_{k+1}))$, as shown in Fig 11. We assign each edge clip (e, l_k) to the sub-tile with the rectangle containing the bounding box of l_k .

Because, by the induction assumption property \mathcal{F} is true on L_z , and by construction, each new edge clip can cross the boundary of the sub-tile only at the clip endpoints. We also cover all the intersections of p with the sub-tiles with the new edge clips, so the property \mathcal{F} holds for L_{z+1} .

Two parameters control the algorithm: tile capacity, \mathcal{C} , and the minimal size of a tile: $(\mathcal{W}, \mathcal{H})$. If for each (i, j) the number of elements in $L_z(i, j).tiledata$ is not greater than \mathcal{C} , and if $w \leq \mathcal{W}$ and $h \leq \mathcal{H}$, where (w, h) is the current tile size, then we try to build the next level L_{z+1} . Otherwise, the second phase starts.

In our setting $\mathcal{C} = 500$, and $(\mathcal{W}, \mathcal{H}) = 3(w, h)$, where w is the average width and h is the average height of the nodes of the graph.

For efficiency, we do not create a new curve in an edge clip but keep two parameters indicating the clipped segment start and end. A possible optimization here is to find the repeated segments in the edge curves that naturally appear while routing through the same graph with the same algorithm, and reuse the repeated segment to save memory and to avoid the same calculation in edge clipping.

4.2 Second phase of tiling

In this phase we filter out some entities from the lower levels. We do not change the highest, the most detailed level. We sort the nodes of the graph into array N




Fig. 11: Intersect curve $[start, end]$ with the midlines. Sort the intersections parameters, together with start and end, into array $u = [start, a, b, c, end]$. Split the curve to sub-curves $[start, a]$, $[a, b]$, $[b, c]$, and $[c, end]$. Each sub-curve is confined to a single sub-tile.

by PageRank [30]. For each level L , except of the highest, we proceed as follows.

```

1: procedure FILTER( $L$ )
2:    $r \leftarrow \text{removeEntities}(L)$ 
3:   for all  $n$  in  $N$  do
4:     if ! $\text{addNodeToLevel}(n, r, N)$  then break
5:   end if
6: end for
7: end procedure

```

Here $\text{removeEntities}(L)$ empties all the tiles of level L , and returns map r allowing to restore the tiles. Map r maps each graph element to an array of tile elements representing it in L . Function $\text{addNodeToLevel}(n)$ tries to add node n to L , it also tries to add the tile elements for self edges of n , and the tile elements for the edges connecting n with the nodes ranked at least as high as n . These nodes are the nodes already added to L .

This procedure guarantees that each tile of L has no more than \mathcal{C} elements.

4.3 Third phase of tiling

In the third phase we use the fact that some nodes are not present on the level. For all levels, except of the highest, we reroute the edges but only around the nodes that are present in the level. We do not calculate edge routes from scratch, but use the existing routes and only apply the "funnel" heuristic in

larger channels. This gives us simpler edge routes but still has the visual stability during the level change while browsing.

5 Conclusion

The first contribution is an efficient edge routing algorithm. The algorithm replaces each polyline path with the shortest path that does not intersect the nodes and stays in the polyline homotopy class, the channel, and then produces the smooth composite curve. It uses a modification of the 'funnel' in the simple polygon algorithm, where we drop the requirement that the polygon is simple.

We described the scenario where the graph spanner is used. Instead, our approach can start with any routing with polylines outside the nodes.

The algorithm is fast and creates visually pleasing results.

The tiling method is the second contribution of this study. It allows large graphs to be visualized in a web browser, similar to online maps. The method provides an overview of the most important nodes and edges of the graph by making them visible on the top levels while keeping the number of visible elements under a given limit. The novelty of the method is that it efficiently subdivides the edges by using the tiles, and simplifies the edges on the upper levels with the edge routing algorithm mentioned above.

6 Future work

- Find a tiling method that guarantees that each tile has no more than \mathcal{C} elements on every level. One approach is to use more aggressive, and regular edge bundling to reduce the number of edge clips in the tiles.
- Our tile calculation is memory intensive and takes a longer time for larger graphs. The largest graph from the Table 1 that we were able to load with Chrome, and Edge using the tiling procedure was p2p-Gnutella04 [26]. One of the reasons was the memory limit on a process in those browsers, another was the long running-time of the tiling procedure. A possible measure would be saving the tiles to the disk and loading them on demand.
- For the user convenience we would like to run the layout, routing, and tiling, in a worker thread to avoid blocking the main thread.
- Addressing node labels visibility is an important task. We would like to enlarge the most important nodes of the view so that their labels are readable.

References

1. “Graphviz.” <http://www.graphviz.org/>.
2. “sfdp.” <https://graphviz.org/docs/layouts/sfdp/>.
3. B. Chazelle, “A theorem on polygon cutting with applications,” in *23rd Annual Symposium on Foundations of Computer Science (sfcs 1982)*, pp. 339–349, IEEE, 1982.

4. J. Hershberger and J. Snoeyink, "Computing minimum length paths of a given homotopy class," *Computational geometry*, vol. 4, no. 2, pp. 63–97, 1994.
5. "yworks." <https://yworks.com/products/yed>.
6. "Regraph." <https://cambridge-intelligence.com/regraph/>.
7. "Skewed." <https://graph-tool.skewed.de>.
8. "Circos." <http://circos.ca/>.
9. "Cosmograph." <https://cosmograph.app>.
10. L. Nachmanson, R. Prutkin, B. Lee, N. H. Riche, A. E. Holroyd, and X. Chen, "Graphmaps: Browsing large graphs as interactive maps," in *Graph Drawing and Network Visualization: 23rd International Symposium, GD 2015, Los Angeles, CA, USA, September 24-26, 2015, Revised Selected Papers 23*, pp. 3–15, Springer, 2015.
11. A. Perrot and D. Auber, "Cornac: Tackling huge graph visualization with big data infrastructure," *IEEE Transactions on Big Data*, vol. 6, no. 1, pp. 80–92, 2018.
12. C. Hurter, O. Ersoy, and A. Telea, "Graph bundling by kernel density estimation," in *Computer graphics forum*, vol. 31, pp. 865–874, Wiley Online Library, 2012.
13. K. Xu, C. Rooney, P. Passmore, D.-H. Ham, and P. H. Nguyen, "A user study on curved edges in graph visualization," *IEEE transactions on visualization and computer graphics*, vol. 18, no. 12, pp. 2449–2456, 2012.
14. D. Holten and J. J. Van Wijk, "A user study on visualizing directed edges in graphs," in *Proceedings of the SIGCHI conference on human factors in computing systems*, pp. 2299–2308, 2009.
15. T. Dwyer and L. Nachmanson, "Fast edge-routing for large graphs," in *Graph Drawing: 17th International Symposium, GD 2009, Chicago, IL, USA, September 22-25, 2009. Revised Papers 17*, pp. 147–158, Springer, 2010.
16. B. Flinchbaugh and L. Jones, "Strong connectivity in directional nearest-neighbor graphs," *SIAM Journal on Algebraic Discrete Methods*, vol. 2, no. 4, pp. 461–463, 1981.
17. A. C.-C. Yao, "On constructing minimum spanning trees in k-dimensional spaces and related problems," *SIAM Journal on Computing*, vol. 11, no. 4, pp. 721–736, 1982.
18. D. P. Dobkin, E. R. Gansner, E. Koutsofios, and S. C. North, "Implementing a general-purpose edge router," in *Graph Drawing: 5th International Symposium, GD'97 Rome, Italy, September 18-20, 1997 Proceedings 5*, pp. 262–271, Springer, 1997.
19. B. Delaunay *et al.*, "Sur la sphere vide," *Izv. Akad. Nauk SSSR, Otdelenie Matematicheskii i Estestvennyka Nauk*, vol. 7, no. 1, pp. 793–800, 1934.
20. V. Domiter and B. Žalik, "Sweep-line algorithm for constrained delaunay triangulation," *International Journal of Geographical Information Science*, vol. 22, no. 4, pp. 449–462, 2008.
21. "Funnel algorithm." <https://page.mi.fu-berlin.de/mulzer/notes/alggeo/polySP.pdf>.
22. A. Beveridge and M. Chemers, "The game of game of thrones: Networked concordances and fractal dramaturgy," in *Reading Contemporary Serial Television Universes*, pp. 201–225, Routledge, 2018.
23. "b103." <https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot>.
24. "b100." <https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot>.
25. "Skewed." <http://mozart.diei.unipg.it/gdcontest/contest2011/composers.xml>.
26. "p2p-gnutella04." <https://snap.stanford.edu/data/p2p-Gnutella04.html>.
27. "facebookcombined." https://snap.stanford.edu/data/facebook_combined.txt.gz.

28. B. Rozemberczki and R. Sarkar, “Characteristic Functions on Graphs: Birds of a Feather, from Statistical Descriptors to Parametric Models,” in *Proceedings of the 29th ACM International Conference on Information and Knowledge Management (CIKM '20)*, p. 1325–1334, ACM, 2020.
29. J. Leskovec, J. Kleinberg, and C. Faloutsos, “Graph evolution: Densification and shrinking diameters,” *ACM transactions on Knowledge Discovery from Data (TKDD)*, vol. 1, no. 1, pp. 2–es, 2007.
30. L. Page, S. Brin, R. Motwani, and T. Winograd, “The pagerank citation ranking: Bringing order to the web,” *Stanford InfoLab*, vol. 249, no. 373, pp. 1–17, 1999.