

Browsing large graphs with MSAGLJS, a graph draph drawing tool in JavaScript

Lev Nachmanson and Xiaoji Chen

Microsoft Research, US,
levnach@hotmail.com, cxiaoji@gmail.com,
Msagljs github home page: <https://github.com/microsoft/msagljs>

1 **Abstract.** There has been progress in visualization of large graphs re-
2 cently. Tools appeared that can render a huge graph in seconds. However,
3 if we request that the node labels are rendered, and the edges are routed
4 around the nodes, then the problem is still standing. Interacting with a
5 large graph in an Internet browser with the same ease as browsing an
6 online map, inspecting the high level structure and zooming in to the
7 high level detail, is still a challenging task.
8 In this paper we describe novel approaches to several aspects of this
9 problem.
10 We give a new efficient algorithm for edge routing, where the edges are
11 routed around the nodes. The algorithm produces edge paths which are
12 visually appealing and optimal in their homotopy class.
13 To facilitate graph visualization with DeckGL, or any other viewer sup-
14 porting tiles, we propose a new simple and fast tiling method. The
15 method guarantees that in every view, except of the highest layer, the
16 number of visible nodes per tile is not larger than a predefined bound.
17 Our method provides a high level overview of the graph, with the grad-
18 ual increase of the detail level. We make the node labels of the most
19 important nodes for the current view visible.
 The edge routing algorithm mentioned above is reused at the tiling stage
 to simplify the paths on the lower levels. In addition, we bundle edges
 per-tile as an optimization heuristic

20 Introduction

21 We target large but not huge graphs. The maximum number of vertices of the
22 graphs we looked at was 28k, and the maximum number of the edges was 237k.
23 There are quite a few algorithms that calculate node positions for such graphs,
24 and work very fast [1, 2]. We look at the node layout as a solved problem.

25 In the first part of the paper we address edge routing where an edge does
26 not intersects the nodes it is not adjacent to. Our approach works for any node
27 layout, as long id does not produce node overlaps. We build on the edge routing
28 from [3] and improve it. There has been progress in visualization of large graphs
29 recently. Tools appeared that can render a huge graph in seconds. However, the

30 situatiton changes if we request that the node labels are rendered, and the edges
31 overlap only the nodes they are adjacent to. Interacting with a large graph in an
32 Internet browser with the same ease as browsing an online map, inspecting the
33 high level structure and zooming in to the high level detail, is still an unsolved
34 problem. In this paper we describe novel approaches to several aspects of this
35 problem.

36 We propose a novel and efficient algorithm for edge routing, where each edge
37 can only intersect its source or target. The algorithm produces edge paths which
38 are visually appealing and even optimal in their homotopy class.

39 To facilitate graph visualization with DeckGL, we propose a new simple and
40 fast tiling method. The method guarantees that in every view, except of the
41 views of the Shighest layer, the number of visible entities is not larger than a
42 predefined bound. The method can be used in other viewers that support tiling.

43 Our method provides a high level overview of the graph.

44 The edge routing algorithm mentioned above is reused at the tiling stage to
45 simplify the paths on the lower levels. In addition, we bundle edges per-tile as
46 an optimization heuristic. Our software runs calculations on the client desktop
47 or a phone, and renders the graph in the browser.

48 Related work

49 A popular graph drawing tool Graphviz [4] applies Scalable Force-Directed Place-
50 ment [5] for large graphs, with no support for tiling. Its edge routing for this case
51 builds the whole visibility graph. This can be very slow because the visibility
52 graph can have $O(n^2)$ edges, where n is the number of the nodes in the graph.
53 Interestingly, the funnel algorithm [6, 7], the last step of our approach, is used
54 in Graphviz for the edge routing in the Sugiyama layout. We are not aware of
55 any tool that integrates Graphviz and uses tiling as well.

56 yWorks [8] has method "Organic edge routing" that produces edge routes
57 around the nodes. We could find only a very general description of the method:
58 "The algorithm is based on a force directed layout paradigm. Nodes act as re-
59 pulsive forces on edges in order to guarantee a certain minimal distance between
60 nodes and edges. Edges tend to contract themselves. Using simulated annealing,
61 this finally leads to edge layouts that are calculated for each edge separately".
62 It seems the algorithm runs in $O(n + m)\log(n + m)$ time, where n is the number
63 of the nodes and m is the number of the edges.

64 ReGraph [9] uses WebGL as the viewing platform. It can render a large graph
65 using straight lines for the edges. The tool does not support tiling, but instead
66 the user interactively opens the node that is a cluster of nodes.

67 "graph-tool.skewed" [10] does not implement its own layout algorithms or
68 edge routing algorithms, but instead provides a nice wrapper around the algo-
69 rithms from other layout tools.

70 Circos [11] visualizes large graphs in a circular layout. It does not support
71 tiles.

72 Cosmograph [12] uses a GPU to calculate the layout of a graph and can
 73 handle a graph with a million nodes. It renders edges as straight lines. It does
 74 not support tiling.

75 The authors of [13] implemented GraphMaps, a tool for large graph visual-
 76 ization. The tool only ran in Windows. The edge were routed as polylines on a
 77 triangulation. The tool supported tiling, but the problem of the limiting number
 78 of visible entities was not solved.

79 In [14] an approach to visualize a huge graph. The method uses tiles and
 80 edge bundling following [15], which is applied at the last moment during the the
 81 graph browsing. The latter calculation is done on the client side, the rest of the
 82 calculations run on several server machines.

83 Edge routing in MSAGLJS

94 The edge routing starts, as in [3], by building a spanner graph, an approximation
 95 of the full visibility graph, and then finding shortest paths on the spanner. The
 96 spanner, see Fig 2, is built on a variation of a Yao graph, which was introduced
 97 independently by Flinchbaugh and Jones [16] and Yao [17]. This kind of graph
 98 is defined by the set of cones with the apices at the vertices. The cones have the
 99 same angle, usually in the form of $\frac{2\pi}{n}$, where n is a natural number. The family
 100 of cones with the apex at a specific vertex partition the plane as illustrated in
 101 Fig. 1. For each cone at most one edge is created connecting the cone apex with
 102 a vertex inside of the cone, so the graph has $O(n)$ edges where n is the number
 103 of vertices.

104
 108 The approach of [3] applies local optimizations to shorten an edge path.
 109 Namely, it tries to shortcut one vertex at a time from the path, as illustrated in
 110 Fig 3. To smoothen a path, it fits Bezier segments into the polyline corners by
 111 using a binary search to find a larger fitting segment, see Fig 4. While analyzing
 112 performance of the edge routing in MSAGLJS, we noticed, that for a graph with
 113 more than 1k nodes these heuristics sometime create a performance bottleneck,
 114 in spite of using R-Trees[18].

115 In addition, when the naive shortcutting of polyline corners fails, the resulting
 116 path might remain not visually appealing, as shown in Fig. 3.

117 We replace these heuristics with a more precise and efficient optimization
 118 described below.

119 Path optimization

120 We finalize edge routes by the “funnel” algorithm [6, 7], routing a path inside a
 121 simple polygon, that is a polygon without holes.

122 An application of the ‘path in a simple polygon’ optimization to edge routing
 123 is not a new idea: the novelty of our work is in how we find the polygon and
 124 how we use it. The authors of Graphvis used the ‘funnel’ algorithm [19], but
 125 only for hierarchical layouts, where a simple polygon, \mathcal{P} , containing the path is



Fig. 1. Yao graph




Fig. 2. Spanner graph is built using the idea of Yao graphs. The dashed curves are the original node boundaries. Each original curve is surrounded by a polygon with some offset to allow the polyline paths smoothing without intersecting the former.

The edge marked by the circles is created because the top vertex is inside of the cone and it is the closest among such vertices to the cone apex. The apex of the cone is the lower vertex of the edge.

MSAGLJS uses cone angle $\frac{\pi}{6}$, so the edges of the spanner can deviate from the optimal direction by this angle. Therefore, the shortest paths on the spanner have length that is at most the optimal shortest length multiplied by $\frac{1}{\cos(\frac{\pi}{6})} \simeq 1.155$.



Fig. 3. Unsuccessful shortcut

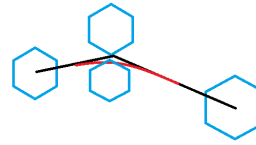





Fig. 4. Fitting a Bezier segment into a polyline corner




139 **Fig. 5.** Path \mathcal{L} with \mathcal{T} , a fragment.



140 **Fig. 6.** Polygon \mathcal{P} containing \mathcal{L} .



141 **Fig. 7.** New triangulation of \mathcal{P} .



142 **Fig. 8.** The optimized path together
143 with the sleeve diagonals.

126 available. They write: "If \mathcal{P} does not contain holes ... we can apply a standard
127 "funnel" algorithm ... for finding Euclidean shortest paths in a simple polygon".
128 In general case, for a non-layered layout, they build the visibility graph which is
129 very expensive for a large graph.

130 Here we find the polygon \mathcal{P} for any layout. We drop the requirement that
131 \mathcal{P} is simple. Indeed, to run the "funnel" algorithm one only needs a "sleeve": a
132 sequence of triangles leading from the start to the end of the path, where each
133 triangle shares a side with its successor. Let us show how to build polygon \mathcal{P} ,
134 create a sleeve, and produce an optimized path.

135 We call obstacles, \mathcal{O} , the set of polygons covering the original nodes, see
136 Fig. 2. Before routing edges, we calculate a Constrained Delaunay Triangulation
137 [20] on \mathcal{O} . Let us call this triangulation \mathcal{T} .

138 For each edge of the graph we proceed with the following steps.


146 We route a path, called \mathcal{L} , on the spanner, as illustrated by Fig. 5. Let \mathcal{S} and
147 \mathcal{E} be the obstacles containing correspondingly \mathcal{L} 's start and end point. To obtain
148 \mathcal{P} , let us consider \mathcal{U} , the set of all triangles $t \in \mathcal{T}$ such that either $t \subset \mathcal{S} \cup \mathcal{E}$, or t
149 intersects \mathcal{L} and is not inside of any obstacle in $\mathcal{O} \setminus \{\mathcal{S}, \mathcal{E}\}$. The union of \mathcal{U} gives
150 us \mathcal{P} . The boundary of \mathcal{P} comprizes all sides e of the triangles from \mathcal{U} such that
151 e belongs to exactly one triangle from \mathcal{U} , see Fig. 6.

152 To create the sleeve [6, 7], we need to have a triangulation of \mathcal{P} such that every
153 edge of the triangulation is either a boundary edge of \mathcal{P} , or a diagonal of \mathcal{P} .
154 Because \mathcal{U} might not have this property, as in Fig. 6, we create a new Constrained
155 Delaunay Triangulation of \mathcal{P} , where the set of constrained edges is the boundary
156 of \mathcal{P} , see Fig. 7.

157 We trace path \mathcal{L} through the new triangulation and obtain the sleeve. Finally,
158 we apply the funnel algorithm on the sleeve and obtain the path which is the
159 shortest in the homotopy class of \mathcal{L} , as illustrated in Fig. 8.

160 The discussion [21] of the algorithm helped us in the implementation.

161 Polygon \mathcal{P} is not necessarily simple, as shown in Fig. 9. In this example the



144 **Fig. 9.** \mathcal{P} is not simple. The dotted path is shorter than the dashed one that
 145 was found by the routing.

162 path that we calculate with the funnel algorithm is not the shortest path inside
 163 of \mathcal{P} .

164 Performance and quality comparison

168 In Fig. 10 we compare the paths generated by the old and the new method. We
 169 can see that the paths produced by the new method have no kinks. We also
 170 know that these paths are the shortest in their 'channels'. Arguably, the new
 171 method produces better paths.

183 Our performance experiments are summarized in Table. 1. We see that the
 184 older approach outperforms the new one on the smaller graphs; those with the
 185 number of nodes under 2000. The new method is faster on the rest of the graphs.
 186 We still prefer to use the new method independently of the graph size since the
 187 slowdown is insignificant, under half of a second in our experiments, but the
 188 quality of the paths is better. On the larger graphs the new method runs faster
 189 and produces better paths, so it is an obvious choice.

190 1 Tiling

191 The algorithm works in three phases. The first phase builds the levels starting
 192 from the lowest level and proceeding to higher and more detailed levels, with
 193 smaller tiles, until no more tile subdivision is required. The second phase filters
 194 out the entities from the layers to satisfy the capacity quota. Finally, the third




Fig. 10. The difference in the paths between the old, on the left, and the new, on the right, paths. The arrows on the left fragment point to the kinks that were removed by the new method.

graph	nodes	edges	old method's time	new time
social network [22]	407	2639	1.0	1.4
b103 [23]	944	2438	1.6	2.0
b100 [24]	1463	5806	5.6	5.785
composers [25]	3405	13832	510.5	17.5
p2p-Gnutella04 [26]	10876	39994	375.4	293.8
facebook_combined [27]	4039	88234	132.2	119.1
lastfm_asia_edges [28]	7626	27807	43.3	41.4
deezer_europe_edges [28]	28283	92753	1596.9	1209.3
ca-HepPh [29]	12008	237010	521.2	495.0

Table 1. Performance comparison with time in seconds.

195 phase simplifies the edge routes to utilize the space freed by the filtered out
196 entities.

197 A tile, in our settings, is a pair $(rect, tiledata)$, where $rect$ is the rectangle of
198 the tile and $tiledata$ is a set of *tile elements* visible in $rect$. A *tile element* could
199 be a node, an edge label, an edge arrowhead, or an *edge clip*. An edge clip is a
200 pair (e, p) , where e is an edge and p is a continuous piece of the edge curve c_e .
201 Sometimes we need several edge clips to trace an edge through a tile.

202 The initial tile, the only tile on level 0, is represented by pair $(0, 0)$. For
203 $z = 1$, there are four tiles: $(0, 0)$, $(0, 1)$, $(1, 0)$, and $(1, 1)$. Each tile (i, j) can be
204 subdivided into four subtiles for level $z + 1$: $(2i, 2j)$, $(2i, 2j + 1)$, $(2i + 1, 2j)$, and
205 $(2i + 1, 2j + 1)$.

206 Each z -level is represented by a map L_z , so $L_z(i, j)$ gives us a specific tile.
207 Empty tiles correspond to undefined $L_z(i, j)$.

208 We use edge clips to represent the edge intersections with the tiles and provide
209 the renderer with the minimal geometry that is sufficient to render a tile. To
210 achieve this we require property \mathcal{F} :

- 211 a) For each tile t , for each curve clip $(e, p) \in t.tiledata$, we have: $p \subset t.rect$
212 and p might cross the boundary of the $t.rect$ only at endpoints of p .
- 213 b) For each edge e we have : the union of all p for all $(e, p) \in t.tiledata$ is
214 equal to $c_e \cap t.rect$.

215 First phase of tiling


216 The first phase starts with $L_0 = \{(0, 0) \rightarrow (rect, tiledata)\}$: and $tiledata$ com-
217 prising curve clips (e, c_e) , for all edges e of the graph, all graph nodes, all edge
218 labels, and all edge arrowheads. We ensure property \mathcal{F} by setting $rect$ to a
219 padded bounding box of the graph, so each edge curve does not intersect the
220 boundary of $rect$.

221 Let us assume that L_z is already constructed and \mathcal{F} holds for its tiles. To
222 build level L_{z+1} we divide each tile $t = L_z(i, j)$ into four subtiles of equal size.
223 For each node, arrowhead, or edge label of $t.tiledata$, if the bounding box of the
224 element intersects the subtile's rectangle then we add the element to the subtile
225 $tiledata$.

226 The edge clip treatment is more involved. Let (e, p) be a curve clip belonging
227 to tile t . We find all intersections of curve p with the horizontal midline and the
228 vertical midline of $t.rect$. Each intersection can be represented as $p[t_j]$. We sort
229 sequence $u = [start, \dots, t_j, \dots, end]$, where $[start, end]$ is the parameter domain
230 of p , in ascending order, and remove the duplicates.

231 Next we create curve clips $(e, l_k) = (e, trim(p, u_k, u_{k+1}))$, as shown in Fig 11.
232 We assign each curve clip (e, l_k) to the subtile with the rectangle containing the
233 bounding box of l_k .

237 Because, by the induction assumption property \mathcal{F} is true on L_z , and by
238 construction, each new curve clip can cross the boundary of the subtile only at
239 the clip endpoints. We also cover all the intersections of p with the subtiles with
240 the new edge clips, so the property \mathcal{F} holds for L_{z+1} .



234 **Fig. 11.** Intersect curve $[start, end]$ with the midlines. Sort the intersections pa-
 235 rameters together with start, and end into array $u = [start, a, b, c, end]$. Split the
 236 curve to sub-curves $[start, a]$, $[a, b]$, $[b, c]$, $[c, end]$.

241 Two parameters control the algorithm: tile capacity, \mathcal{C} , and the minimal size
 242 of a tile: $(\mathcal{W}, \mathcal{H})$. If for each (i, j) the number of elements in $L_z(i, j).tiledata$ is
 243 not greater than \mathcal{C} , or, if $w \leq \mathcal{W}$ and $h \leq \mathcal{H}$, where w (h) is the current tile
 244 width (correspondingly, height), then the second phase starts.

245 In our setting $\mathcal{C} = 500$, and $(\mathcal{W}, \mathcal{H}) = 3(w, h)$, where w is the average width
 246 and h is the average height of the nodes of the graph.

247 **Edge bundling** In our settings each edge clip is uniquely defined, module
 248 direction, by its start and end point. We can use this property to bundle the
 249 edges. In each tile we keep a map from unordered pairs of points to the set of
 250 edge clips that have these points as start and end points. Each such pair defines
 251 an edge bundle. For all edge clips in a bundle we create only one curve segment,
 252 avoiding the expensive trimming. We also count a bundle as one element in the
 253 tile, as in most of the cases the drawing attributes of the edges in the bundle are
 254 the same.

255 In our experiments, the number of edge bundles is about 50% of the number
 256 of edge clips, so the edge bundling is a significant optimization.

257 Second phase of tiling

258 The second phase of tiling filters out the entities from the lower layers. We do not
 259 change the highest, the most detailed layer. We sort the nodes of the graph into
 260 array N by PageRank [30]. For each layer L , except of the highest, we proceed
 261 as follows.

262 Here `removeEntities(L)` empties all the tiles of layer L , but returns map r allow-
 263 ing to restore the tiles. Function `addNodeToLayer(n)` returns false and does not
 264 change L when one of the tiles intersecting n already has more elements than

```

1: procedure FILTER( $L$ )
2:    $r \leftarrow \text{removeEntities}(L)$ 
3:   for all  $n$  in  $N$  do
4:     if ! $\text{addNodeToLayer}(n, r, N)$  then break
5:   end if
6: end for
7: end procedure

```

265 \mathcal{C} . Otherwise, the function adds n to all tiles intersected by n . It also adds the
266 tile elements for self edges of n , and the tile elements for the edges connecting n
267 with the nodes appearing in N before n , i.e. the nodes with the rank not lesser
268 than the rank of n .

269 This procedure guarantees that each tile of L has no more than \mathcal{C} nodes, but
270 a tile can have more than \mathcal{C} elements in general.

271 Third phase of tiling

272 In the third phase we use a fact that some nodes are not present on the layer. For
273 all layers, except of the highest, we reroute the edges but only around the nodes
274 that are present in the layer. We do not calculate edge routes from scratch, but
275 use the existing routes and only apply the "funnel" heuristic in larger channels.
276 This gives us simpler edge routes but still has a visual stability during the layer
277 change while browsing.

278 2 Future work

- 279 – Find a tiling method that guarantees that each tile has no more than \mathcal{C} el-
280 ements. One approach could be to use a more aggressive edge bundling to
281 reduce the number of edge clips in the tiles.

282 References

- 283 1. Y. Hu and L. Shi, "Visualizing large graphs," *Wiley Interdisciplinary Reviews:*
284 *Computational Statistics*, vol. 7, no. 2, pp. 115–136, 2015.
- 285 2. U. Brandes and C. Pich, "Eigensolver methods for progressive multidimensional
286 scaling of large data," in *Graph Drawing: 14th International Symposium, GD*
287 *2006, Karlsruhe, Germany, September 18-20, 2006. Revised Papers 14*, pp. 42–
288 53, Springer, 2007.
- 289 3. T. Dwyer and L. Nachmanson, "Fast edge-routing for large graphs," in *Graph*
290 *Drawing: 17th International Symposium, GD 2009, Chicago, IL, USA, September*
291 *22-25, 2009. Revised Papers 17*, pp. 147–158, Springer, 2010.
- 292 4. "Graphviz." <http://www.graphviz.org/>.
- 293 5. "sfdp." <https://graphviz.org/docs/layouts/sfdp/>.

- 294 6. B. Chazelle, "A theorem on polygon cutting with applications," in *23rd Annual*
295 *Symposium on Foundations of Computer Science (sfcs 1982)*, pp. 339–349, IEEE,
296 1982.
- 297 7. J. Hershberger and J. Snoeyink, "Computing minimum length paths of a given
298 homotopy class," *Computational geometry*, vol. 4, no. 2, pp. 63–97, 1994.
- 299 8. "yworks." <https://yworks.com/products/yed>.
- 300 9. "Regraph." <https://cambridge-intelligence.com/regraph/>.
- 301 10. "Skewed." <https://graph-tool.skewed.de>.
- 302 11. "Circos." <http://circos.ca/>.
- 303 12. "Cosmograph." <https://cosmograph.app>.
- 304 13. L. Nachmanson, R. Prutkin, B. Lee, N. H. Riche, A. E. Holroyd, and X. Chen,
305 "Graphmaps: Browsing large graphs as interactive maps," in *Graph Drawing and*
306 *Network Visualization: 23rd International Symposium, GD 2015, Los Angeles, CA,*
307 *USA, September 24–26, 2015, Revised Selected Papers 23*, pp. 3–15, Springer, 2015.
- 308 14. A. Perrot and D. Auber, "Cornac: Tackling huge graph visualization with big data
309 infrastructure," *IEEE Transactions on Big Data*, vol. 6, no. 1, pp. 80–92, 2018.
- 310 15. C. Hurter, O. Ersoy, and A. Telea, "Graph bundling by kernel density estimation,"
311 in *Computer graphics forum*, vol. 31, pp. 865–874, Wiley Online Library, 2012.
- 312 16. B. Flinchbaugh and L. Jones, "Strong connectivity in directional nearest-neighbor
313 graphs," *SIAM Journal on Algebraic Discrete Methods*, vol. 2, no. 4, pp. 461–463,
314 1981.
- 315 17. A. C.-C. Yao, "On constructing minimum spanning trees in k-dimensional spaces
316 and related problems," *SIAM Journal on Computing*, vol. 11, no. 4, pp. 721–736,
317 1982.
- 318 18. A. Guttman, "R-trees: A dynamic index structure for spatial searching," in *Pro-*
319 *ceedings of the 1984 ACM SIGMOD international conference on Management of*
320 *data*, pp. 47–57, 1984.
- 321 19. D. P. Dobkin, E. R. Gansner, E. Koutsofios, and S. C. North, "Implementing a
322 general-purpose edge router," in *Graph Drawing: 5th International Symposium,*
323 *GD'97 Rome, Italy, September 18–20, 1997 Proceedings 5*, pp. 262–271, Springer,
324 1997.
- 325 20. B. Delaunay, "Sur la sphere vide, bull. acad. science ussr vii: Class," *Sci. Mat. Nat*,
326 pp. 793–800, 1934.
- 327 21. "Funnel algorithm." <https://page.mi.fu-berlin.de/mulzer/notes/alggeo/polySP.pdf>.
- 328 22. A. Beveridge and M. Chemers, "The game of game of thrones: Networked con-
329 cordances and fractal dramaturgy," in *Reading Contemporary Serial Television*
330 *Universes*, pp. 201–225, Routledge, 2018.
- 331 23. "b103." [https://github.com/microsoft/automatic-graph-](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot)
332 [layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot).
- 333 24. "b100." [https://github.com/microsoft/automatic-graph-](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot)
334 [layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot).
- 335 25. "Skewed." <http://mozart.diei.unipg.it/gdcontest/contest2011/composers.xml>.
- 336 26. "p2p-gnutella04." <https://snap.stanford.edu/data/p2p-Gnutella04.html>.
- 337 27. "facebookcombined." <https://snap.stanford.edu/data/facebookcombined.txt.gz>.
- 338 28. B. Rozemberczki and R. Sarkar, "Characteristic Functions on Graphs: Birds of a
339 Feather, from Statistical Descriptors to Parametric Models," in *Proceedings of the*
340 *29th ACM International Conference on Information and Knowledge Management*
341 *(CIKM '20)*, p. 1325–1334, ACM, 2020.
- 342 29. J. Leskovec, J. Kleinberg, and C. Faloutsos, "Graph evolution: Densification
343 and shrinking diameters," *ACM transactions on Knowledge Discovery from Data*
344 *(TKDD)*, vol. 1, no. 1, pp. 2–es, 2007.

- 345 30. L. Page, S. Brin, R. Motwani, and T. Winograd, “The pagerank citation ranking:
346 Bringing order to the web,” *Stanford InfoLab*, vol. 249, no. 373, pp. 1–17, 1999.