

Browsing large graphs with XJS, a graph drawing tool in JavaScript

auth0 and auth1

Institute, US,
auth0@hotmail.com, auth1@gmail.com,
X github home page: <https://github.com/X>

1 **Abstract.** There has been progress in visualization of large graphs re-
2 cently. Tools appeared that can render a huge graph in seconds. However,
3 if we request that the node labels are readable, and the edges are routed
4 around the nodes, then the problem remains difficult. Interacting with a
5 large graph in a web browser with the same ease as browsing an online
6 map is still a challenging task. In this paper we describe a few novel ap-
7 proaches to large graph visualization that we developed in open-source
8 JavaScript software.
9 We give a new efficient edge routing algorithm, where the edges are
10 routed around the nodes. The algorithm produces edge paths which are
11 visually appealing and shortest in their homotopy class.
12 To facilitate graph visualization with WebGL, or any other platform
13 supporting tiles, we propose a new simple and efficient tiling method.
14 The method guarantees that in every view, except of the highest level,
15 the number of visible entities per tile is not larger than a predefined
16 bound.
 The edge routing algorithm mentioned above is reused at the tiling stage
 to simplify the paths on the upper levels.

17 Introduction

18 Our software is open source written in TypeScript, it is consumed as a set of
19 NPM packages. It runs on the client desktop or on a phone, and renders the graph
20 in a web browser. We target large but not huge graphs. The maximum number of
21 vertices of the graphs we applied our tool to was 28k, and the maximum number
22 of the edges was 237k.

23 The methods and algorithms described below are implemented in XJS.

24 The rest of the paper is divided into sections, including Related Work, Edge
25 routing, Tiling, Conclusion, and Future work.

26 Let us start with a short review of some relevant to us publications.

27 Related work

28 A popular graph drawing tool Graphviz [1] applies method Scalable Force-
29 Directed Placement [2] for large graphs, with no support for tiling. The edge

30 routing for this method builds the whole visibility graph and routes edges on it.
 31 This can be very slow because the visibility graph can have $O(n^2)$ edges, where
 32 n is the number of the nodes in the graph. Interestingly, the funnel algorithm [3,
 33 4], the last step of our approach, is used in Graphviz for the edge routing in the
 34 Sugiyama layout. We are not aware of any tool that integrates Graphviz and
 35 uses tiling as well.

36 yWorks [5] has method "Organic edge routing" that produces edge routes
 37 around the nodes. We could find only a very general description of the method:
 38 "The algorithm is based on a force directed layout paradigm. Nodes act as re-
 39 pulsive forces on edges in order to guarantee a certain minimal distance between
 40 nodes and edges. Edges tend to contract themselves. Using simulated annealing,
 41 this finally leads to edge layouts that are calculated for each edge separately".
 42 It seems the algorithm runs in $O(n+m)\log(n+m)$ time, where n is the number
 43 of the nodes and m is the number of the edges.

44 ReGraph [6] uses WebGL as the viewing platform. It can render a large graph
 45 using straight lines for the edges. The tool does not support tiling, but instead
 46 the user interactively opens the node that is a cluster of nodes.

47 "graph-tool.skewed" [7] does not implement its own layout algorithms or edge
 48 routing algorithms, but instead provides a nice wrapper around the algorithms
 49 from other layout tools.

50 Circos [8] visualizes large graphs in a circular layout. It does not support
 51 tiles.

52 Cosmograph [9] uses a GPU to calculate the layout of a graph and can
 53 handle a graph with a million nodes. It renders edges as straight lines. It does
 54 not support tiling.

55 The authors of [10] implemented GraphMaps, a tool for large graph visu-
 56 alization. The tool only runs on Windows. The edge were routed as polylines
 57 on a triangulation and were not optimized. The tool supported tiling, but the
 58 problem of the limiting number of visible entities was not solved.

59 In [11] an approach to visualize a huge graph is described. The method uses
 60 tiles and edge bundling following [12], which is applied at the last moment during
 61 the graph browsing. The latter calculation is done on the client side. The rest
 62 and the majority of the calculations runs on several servers.

63 Edge routing

64 The user study of Xu et al [13] shows that straight edges improve the user
 65 comprehension of a graph drawing. From the other hand, Holten et al [14] show
 66 that strongly curved edges do not perform well in this regard. In our routing, we
 67 try to keep the edges as straight as possible, and to curve them just enough to
 68 avoid the nodes. Our motivation for developing the routing described here was
 69 to improve the quality of edges and the algorithm speed.

80 The edge routing starts, as in [15], by building a spanner graph, an approxi-
 81 mation of the full visibility graph, and then finding the edge routes as shortest
 82 paths on the spanner. The spanner, see Fig 2, is built on a variation of a Yao

graph, which was introduced independently by Flinchbaugh, and Jones [16], and Yao [17]. This graph is built with a help of a set of cones with the apices at the vertices. Each cone of the set has the same angle, usually in the form of $\frac{2\pi}{k}$, where k is a natural number, $k = 12$ in our settings. The family of cones with the apex at a specific vertex partition the plane, as illustrated in Fig. 1. For each cone at most one edge is created connecting the cone apex with a vertex inside the cone. This way the spanner has at most kn edges, where n is the number of the vertices. We cover each node by a polygon with a relatively small number of corners, at most 8. Polygon corners play role of the vertices of the spanner. As a result, the spanner has $O(N)$ edges, where N is the number of the graph nodes.




Fig. 1. Yao graph




Fig. 2. Spanner graph is built using the idea of Yao graphs. The dashed curves are the original node boundaries. Each original curve is surrounded by a polygon with some offset to allow the polyline paths smoothing without intersecting the former. The edge marked by the circles is created because the top vertex is inside the cone, and it is the closest among such vertices to the cone apex. The apex of the cone is the lower vertex of the edge. XJS uses cone angle $\frac{\pi}{6}$, so the edges of the spanner can deviate from the optimal direction by this angle. Therefore, the shortest paths on the spanner have length that is at most the optimal shortest length multiplied by $\frac{1}{\cos(\frac{\pi}{6})} \simeq 1.155$.




Fig. 3. Unsuccessful shortcut




Fig. 4. Fitting a Bezier segment into a polyline corner

The approach of [15] applies local optimizations to shorten an edge path. Namely, it tries to shortcut one vertex at a time from the path, as illustrated in Fig 3. To smoothen a path, it fits Bezier segments into the polyline corners by using a binary search to find a large fitting segment, see Fig 4. We noticed that when the shortcutting of polyline corners fails, the resulting path might remain not visually appealing, as shown in Fig. 3. We replace the shortcutting with a more precise, but still efficient optimization described below: that is one of the main contributions of our work.

Path optimization

We finalize edge routes by a slight modification of the “funnel” algorithm [3, 4], routing a path inside a simple polygon, that is a polygon without holes.

An application of the ‘path in a simple polygon’ optimization to edge routing is not a new idea: the novelty of our work is in how we find the polygon and how we use it. The authors of Graphvis used the ‘funnel’ algorithm [18], but only for hierarchical layouts, where a simple polygon, \mathcal{P} , containing the path is available. They write: “If \mathcal{P} does not contain holes ... we can apply a standard “funnel” algorithm ... for finding Euclidean shortest paths in a simple polygon”. In general case they build the visibility graph which is very expensive for a large graph.

Here we find the polygon \mathcal{P} for any layout. We drop the requirement that \mathcal{P} is simple. Indeed, to run the “funnel” algorithm one only needs a “sleeve”: a sequence of triangles leading from the start to the end of the path, where each triangle shares a side with its successor. Let us show how to build polygon \mathcal{P} , create a sleeve, and produce an optimized path.

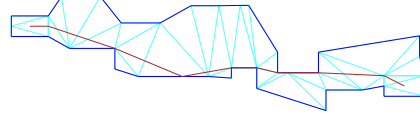
We call obstacles, \mathcal{O} , the set of polygons covering the original nodes, see Fig. 2. Before routing edges, we calculate a Constrained Delaunay Triangulation [19] on \mathcal{O} , following [20]. Let us call this triangulation \mathcal{T} .

For each edge of the graph we proceed with the following steps.

We route a path, called \mathcal{L} , on the spanner, as illustrated by Fig. 5. Let \mathcal{S} and \mathcal{E} be the obstacles containing, correspondingly, \mathcal{L} ’s start and end point. To obtain \mathcal{P} , let us consider \mathcal{U} , the set of all triangles $t \in \mathcal{T}$ such that either



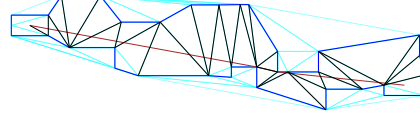
124 **Fig. 5.** Path \mathcal{L} with \mathcal{T} , a fragment.



125 **Fig. 6.** Polygon \mathcal{P} containing \mathcal{L} .



126 **Fig. 7.** New triangulation of \mathcal{P} .



127 **Fig. 8.** The optimized path together
128 with the sleeve diagonals.

134 $t \subset \mathcal{S} \cup \mathcal{E}$, or t intersects \mathcal{L} and is not inside any obstacle in $\mathcal{O} \setminus \{S, E\}$. The
135 union of \mathcal{U} gives us \mathcal{P} . The boundary of \mathcal{P} comprizes all sides e of the triangles
136 from \mathcal{U} such that e belongs to exactly one triangle from \mathcal{U} , see Fig. 6.

137 To create the sleeve [3, 4], we need to have a triangulation of \mathcal{P} such that every
138 edge of the triangulation is either a boundary edge of \mathcal{P} , or a diagonal of \mathcal{P} .
139 Because \mathcal{U} might not have this property, as in Fig. 6, we create a new Constrained
140 Triangulation of \mathcal{P} , where the set of constrained edges is the boundary of \mathcal{P} , see
141 Fig. 7.

142 We trace path \mathcal{L} through the new triangulation and obtain the sleeve. Finally,
143 we apply the funnel algorithm on the sleeve and obtain the path which is the
144 shortest in the homotopy class of \mathcal{L} , as illustrated in Fig. 8.

145 The discussion [21] of the algorithm helped us in the implementation of the
146 funnel algorithm.

147 Polygon \mathcal{P} is not necessarily simple, as shown in Fig. 9. In this example the
148 path that we calculate with the funnel algorithm is not the shortest path inside
149 \mathcal{P} .

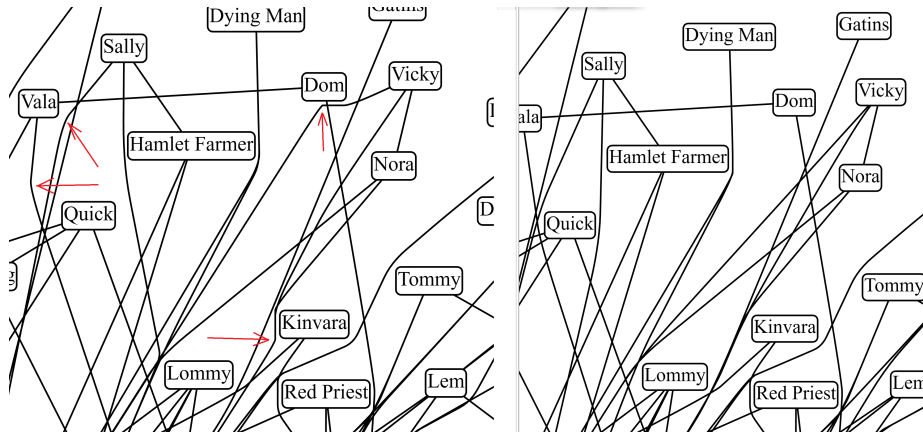
150 Performance and quality comparison

154 In Fig. 10 we compare the paths generated by the old and the new method. We
155 can see that the paths produced by the new method have no kinks. We also know
156 that these paths are the shortest in their 'channels'. Arguably, the new method
157 produces better paths.

158 Our performance experiments are summarized in Table. 1. We see that the
159 older approach outperforms the new one on the smaller graphs; those with the
160 number of nodes under 2000. The new method is faster on the rest of the graphs.
161 We prefer the new method regardless of the graph size because it provides better
162 path quality and the slowdown is insignificant.



129 **Fig. 9.** \mathcal{P} is not simple. The dotted path is shorter than the dashed one that
 130 was found by the routing.



151 **Fig. 10.** Comparing the old, on the left, and the new, on the right, paths. The
 152 arrows on the left fragment point to the kinks that were removed by the new
 153 method.

graph	nodes	edges	old method's time	new time
social network [22]	407	2639	1.0	1.4
b103 [23]	944	2438	1.6	2.0
b100 [24]	1463	5806	5.6	5.785
composers [25]	3405	13832	510.5	20.3
p2p-Gnutella04 [26]	10876	39994	375.4	304.2
facebook_combined [27]	4039	88234	132.2	123.7
lastfm_asia_edges [28]	7626	27807	43.3	54.7
deezer_europe_edges [28]	28283	92753	1596.9	1402.6
ca-HepPh [29]	12008	237010	521.2	495.0

Table 1. Performance comparison with time in seconds.

1 Tiling

We had two goals when working on tiling. The first goal was to make exploring the graph in our tool similar to using online maps. The second goal was efficiency. The algorithm works in three phases. The first phase builds the levels starting from the lowest level and proceeding to higher and more detailed levels, with smaller tiles, until no more tile subdivision is required. The second phase filters out the entities from the layers to satisfy the capacity quota, as in [10]. Finally, the third phase simplifies the edge routes to utilize the space freed by the filtered out entities.

A tile, in our settings, is a pair $(rect, tiledata)$, where $rect$ is the rectangle of the tile and $tiledata$ is a set of *tile elements* visible in $rect$. A *tile element* could be a node, an edge label, an edge arrowhead, or an *edge clip*. An edge clip is a pair (e, p) , where e is an edge and p is a continuous piece of the edge curve c_e . Sometimes we need several edge clips to trace an edge through a tile.


The initial tile, the only tile on level 0, is represented by pair $(0, 0)$. For $z = 1$, there are four tiles: $(0, 0)$, $(0, 1)$, $(1, 0)$, and $(1, 1)$. Each tile (i, j) can be subdivided into four sub-tiles for level $z + 1$: $(2i, 2j)$, $(2i, 2j + 1)$, $(2i + 1, 2j)$, and $(2i + 1, 2j + 1)$.

Each z -level is represented by a map L_z , so $L_z(i, j)$ gives us a specific tile. Empty tiles correspond to undefined $L_z(i, j)$.

We use edge clips to represent the edge intersections with the tiles and provide the renderer with the minimal geometry that is sufficient to render a tile. To achieve this we require property \mathcal{F} :

a) For each tile t , for each edge clip $(e, p) \in t.tiledata$, we have: $p \subset t.rect$ and p might cross the boundary of the $t.rect$ only at endpoints of p .

b) For each edge e we have : the union of all p for all $(e, p) \in t.tiledata$ is equal to $c_e \cap t.rect$.



220 **Fig. 11.** Intersect curve $[start, end]$ with the midlines. Sort the intersections pa-
 221 rameters, together with start and end, into array $u = [start, a, b, c, end]$. Split the
 222 curve to sub-curves $[start, a]$, $[a, b]$, $[b, c]$, and $[c, end]$. Each sub-curve is confined
 223 to a single sub-tile.

201 First phase of tiling

202 The first phase starts with $L_0 = \{(0, 0) \rightarrow (rect, tiledata)\}$: and *tiledata* compris-
 203 ing edge clips (e, c_e) , for all edges e of the graph, all graph nodes, all edge labels,
 204 and all edge arrowheads. We ensure property \mathcal{F} by setting *rect* to a padded
 205 bounding box of the graph, so each edge curve does not intersect the boundary
 206 of *rect*.

207 Let us assume that L_z is already constructed and \mathcal{F} holds for its tiles. To
 208 build level L_{z+1} we divide each tile $t = L_z(i, j)$ into four sub-tiles of equal size.
 209 For each node, arrowhead, or edge label of $t.tiledata$, if the bounding box of the
 210 element intersects the sub-tile's rectangle then we add the element to the sub-tile
 211 *tiledata*.

212 The edge clip treatment is more involved. Let (e, p) be an edge clip belonging
 213 to tile t . We find all intersections of curve p with the horizontal midline and the
 214 vertical midline of $t.rect$. Each intersection can be represented as $p[t_j]$. We sort
 215 sequence $u = [start, \dots, t_j, \dots, end]$, where $[start, end]$ is the parameter domain
 216 of p , in ascending order, and remove the duplicates.

217 Next we create edge clips $(e, l_k) = (e, trim(p, u_k, u_{k+1}))$, as shown in Fig 11.
 218 We assign each edge clip (e, l_k) to the sub-tile with the rectangle containing the
 219 bounding box of l_k .

224 Because, by the induction assumption property \mathcal{F} is true on L_z , and by
 225 construction, each new edge clip can cross the boundary of the sub-tile only at
 226 the clip endpoints. We also cover all the intersections of p with the sub-tiles with
 227 the new edge clips, so the property \mathcal{F} holds for L_{z+1} .

228 Two parameters control the algorithm: tile capacity, \mathcal{C} , and the minimal size
 229 of a tile: $(\mathcal{W}, \mathcal{H})$. If for each (i, j) the number of elements in $L_z(i, j).tiledata$
 230 is not greater than \mathcal{C} , and if $w \leq \mathcal{W}$ and $h \leq \mathcal{H}$, where (w, h) is the current

231 tile size, then we try to build the next level L_{z+1} . Otherwise, the second phase
 232 starts.

233 In our setting $\mathcal{C} = 500$, and $(\mathcal{W}, \mathcal{H}) = 3(w, h)$, where w is the average width
 234 and h is the average height of the nodes of the graph.

235 For efficiency, we do not create a new curve in an edge clip but keep two
 236 parameters indicating the clipped segment start and end. A possible optimization
 237 here is to find the repeated segments in the edge curves that naturally appear
 238 while routing through the same graph with the same algorithm, and reuse the
 239 repeated segment to save memory and to avoid the same calculation in edge
 240 clipping.

241 Second phase of tiling

242 In this phase we filter out some entities from the lower levels. We do not change
 243 the highest, the most detailed level. We sort the nodes of the graph into array N
 244 by PageRank [30]. For each level L , except of the highest, we proceed as follows.

```

1: procedure FILTER( $L$ )
2:    $r \leftarrow \text{removeEntities}(L)$ 
3:   for all  $n$  in  $N$  do
4:     if ! $\text{addNodeToLevel}(n, r, N)$  then break
5:   end if
6: end for
7: end procedure

```

245 Here $\text{removeEntities}(L)$ empties all the tiles of level L , and returns map r al-
 246 lowing to restore the tiles. Map r maps each graph element to an array of tile
 247 elements representing it in L . Function $\text{addNodeToLevel}(n)$ tries to add node n
 248 to L , it also tries to add the tile elements for self edges of n , and the tile elements
 249 for the edges connecting n with the nodes ranked at least as high as n . These
 250 nodes are the nodes already added to L .

251 This procedure guarantees that each tile of L has no more than \mathcal{C} elements.

253 Third phase of tiling

254 In the third phase we use the fact that some nodes are not present on the
 255 level. For all levels, except of the highest, we reroute the edges but only around
 256 the nodes that are present in the level. We do not calculate edge routes from
 257 scratch, but use the existing routes and only apply the "funnel" heuristic in
 258 larger channels. This gives us simpler edge routes but still has the visual stability
 259 during the level change while browsing.

260 2 Conclusion

261 The first contribution is an efficient edge routing algorithm. The algorithm re-
262 places each polyline path with the shortest path that does not intersect the
263 nodes and stays in the polyline homotopy class, the channel, and then produces
264 the smooth composite curve. It uses a modification of the 'funnel' in the simple
265 polygon algorithm, where we drop the requirement that the polygon is simple.

266 We described the scenario where the graph spanner is used. Instead, our
267 approach can start with any routing with polylines outside the nodes.

268 The algorithm is fast and creates visually pleasing results.

269 The tiling method is the second contribution of this study. It allows large
270 graphs to be visualized in a web browser, similar to online maps. The method
271 provides an overview of the most important nodes and edges of the graph by
272 making them visible on the top levels while keeping the number of visible el-
273 ements under a given limit. The novelty of the method is that it efficiently
274 subdivides the edges by using the tiles, and simplifies the edges on the upper
275 levels with the edge routing algorithm mentioned above.

276 3 Future work

- 277 – Find a tiling method that guarantees that each tile has no more than \mathcal{C} el-
278 ements on every level. One approach is to use more aggressive, and regular
279 edge bundling to reduce the number of edge clips in the tiles.
- 280 – Our tile calculation is memory intensive and takes a longer time for larger
281 graphs. The largest graph from the Table 1 that we were able to load with
282 Chrome, and Edge using the tiling procedure was p2p-Gnutella04 [26]. One
283 of the reasons was the memory limit on a process in those browsers, another
284 was the long running-time of the tiling procedure. A possible measure would
285 be saving the tiles to the disk and loading them on demand.
- 286 – For the user convenience we would like to run the layout, routing, and tiling,
287 in a worker thread to avoid blocking the main thread.
- 288 – Addressing node labels visibility is an important task. We would like to en-
289 large the most important nodes of the view so that their labels are readable.

290 References

- 291 1. “Graphviz.” <http://www.graphviz.org/>.
- 292 2. “sfdp.” <https://graphviz.org/docs/layouts/sfdp/>.
- 293 3. B. Chazelle, “A theorem on polygon cutting with applications,” in *23rd Annual*
294 *Symposium on Foundations of Computer Science (sfcs 1982)*, pp. 339–349, IEEE,
295 1982.
- 296 4. J. Hershberger and J. Snoeyink, “Computing minimum length paths of a given
297 homotopy class,” *Computational geometry*, vol. 4, no. 2, pp. 63–97, 1994.
- 298 5. “yworks.” <https://yworks.com/products/yed>.
- 299 6. “Regraph.” <https://cambridge-intelligence.com/regraph/>.

- 300 7. “Skewed.” <https://graph-tool.skewed.de>.
- 301 8. “Circos.” <http://circos.ca/>.
- 302 9. “Cosmograph.” <https://cosmograph.app>.
- 303 10. L. Nachmanson, R. Prutkin, B. Lee, N. H. Riche, A. E. Holroyd, and X. Chen,
304 “Graphmaps: Browsing large graphs as interactive maps,” in *Graph Drawing and*
305 *Network Visualization: 23rd International Symposium, GD 2015, Los Angeles, CA,*
306 *USA, September 24-26, 2015, Revised Selected Papers 23*, pp. 3–15, Springer, 2015.
- 307 11. A. Perrot and D. Auber, “Cornac: Tackling huge graph visualization with big data
308 infrastructure,” *IEEE Transactions on Big Data*, vol. 6, no. 1, pp. 80–92, 2018.
- 309 12. C. Hurter, O. Ersoy, and A. Telea, “Graph bundling by kernel density estimation,”
310 in *Computer graphics forum*, vol. 31, pp. 865–874, Wiley Online Library, 2012.
- 311 13. K. Xu, C. Rooney, P. Passmore, D.-H. Ham, and P. H. Nguyen, “A user study
312 on curved edges in graph visualization,” *IEEE transactions on visualization and*
313 *computer graphics*, vol. 18, no. 12, pp. 2449–2456, 2012.
- 314 14. D. Holten and J. J. Van Wijk, “A user study on visualizing directed edges in
315 graphs,” in *Proceedings of the SIGCHI conference on human factors in computing*
316 *systems*, pp. 2299–2308, 2009.
- 317 15. T. Dwyer and L. Nachmanson, “Fast edge-routing for large graphs,” in *Graph*
318 *Drawing: 17th International Symposium, GD 2009, Chicago, IL, USA, September*
319 *22-25, 2009. Revised Papers 17*, pp. 147–158, Springer, 2010.
- 320 16. B. Flinchbaugh and L. Jones, “Strong connectivity in directional nearest-neighbor
321 graphs,” *SIAM Journal on Algebraic Discrete Methods*, vol. 2, no. 4, pp. 461–463,
322 1981.
- 323 17. A. C.-C. Yao, “On constructing minimum spanning trees in k-dimensional spaces
324 and related problems,” *SIAM Journal on Computing*, vol. 11, no. 4, pp. 721–736,
325 1982.
- 326 18. D. P. Dobkin, E. R. Gansner, E. Koutsofios, and S. C. North, “Implementing a
327 general-purpose edge router,” in *Graph Drawing: 5th International Symposium,*
328 *GD’97 Rome, Italy, September 18–20, 1997 Proceedings 5*, pp. 262–271, Springer,
329 1997.
- 330 19. B. Delaunay *et al.*, “Sur la sphere vide,” *Izv. Akad. Nauk SSSR, Otdelenie Matem-*
331 *aticheskii i Estestvennyka Nauk*, vol. 7, no. 1, pp. 793–800, 1934.
- 332 20. V. Domiter and B. Žalik, “Sweep-line algorithm for constrained delaunay triangu-
333 lation,” *International Journal of Geographical Information Science*, vol. 22, no. 4,
334 pp. 449–462, 2008.
- 335 21. “Funnel algorithm.” <https://page.mi.fu-berlin.de/mulzer/notes/alggeo/polySP.pdf>.
- 336 22. A. Beveridge and M. Chemers, “The game of game of thrones: Networked con-
337 cordances and fractal dramaturgy,” in *Reading Contemporary Serial Television*
338 *Universes*, pp. 201–225, Routledge, 2018.
- 339 23. “b103.” [https://github.com/microsoft/automatic-graph-](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot)
340 [layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot).
- 341 24. “b100.” [https://github.com/microsoft/automatic-graph-](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot)
342 [layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot).
- 343 25. “Skewed.” <http://mozart.diei.unipg.it/gdcontest/contest2011/composers.xml>.
- 344 26. “p2p-gnutella04.” <https://snap.stanford.edu/data/p2p-Gnutella04.html>.
- 345 27. “facebookcombined.” https://snap.stanford.edu/data/facebook_combined.txt.gz.
- 346 28. B. Rozemberczki and R. Sarkar, “Characteristic Functions on Graphs: Birds of a
347 Feather, from Statistical Descriptors to Parametric Models,” in *Proceedings of the*
348 *29th ACM International Conference on Information and Knowledge Management*
349 *(CIKM ’20)*, p. 1325–1334, ACM, 2020.

- 350 29. J. Leskovec, J. Kleinberg, and C. Faloutsos, “Graph evolution: Densification
351 and shrinking diameters,” *ACM transactions on Knowledge Discovery from Data*
352 (*TKDD*), vol. 1, no. 1, pp. 2–es, 2007.
- 353 30. L. Page, S. Brin, R. Motwani, and T. Winograd, “The pagerank citation ranking:
354 Bringing order to the web,” *Stanford InfoLab*, vol. 249, no. 373, pp. 1–17, 1999.