

Browsing large graphs with XJS, a graph drawing tool in JavaScript

auth0 and auth1

Institute, US,
 auth0@hotmail.com, auth1@gmail.com,
 X github home page: <https://github.com/X>

1 **Abstract.** There has been progress in visualization of large graphs re-
 2 cently. Tools appeared that can render a huge graph in seconds. However,
 3 if we request that the node labels are visible, and the edges are routed
 4 around the nodes, then the problem remains difficult. Interacting with
 5 a large graph in an Internet browser with the same ease as browsing an
 6 online map is still a challenging task. In this paper we describe a few
 7 novel approaches to large graph visualization that we developed in an
 8 open-source JavaScript software.
 9 We give a new efficient edge routing algorithm, where the edges are
 10 routed around the nodes. The algorithm produces edge paths which are
 11 visually appealing and shortest in their homotopy class.
 12 To facilitate graph visualization with WebGL, or any other platform
 13 supporting tiles, we propose a new simple and efficient tiling method.
 14 The method guarantees that in every view, except of the highest level,
 15 the number of visible entities per tile is not larger than a predefined
 16 bound.
 17 We make the node labels of the most important nodes of the current
 18 view visible.
 The edge routing algorithm mentioned above is reused at the tiling stage
 to simplify the paths on the lower levels.

19 Introduction

20 We target large but not huge graphs. The maximum number of vertices of the
 21 graphs we looked at was 28k, and the maximum number of the edges was 237k.
 22 There are quite a few algorithms that calculate node positions for such graphs,
 23 and work very fast [1, 2]. We look at the node layout as a solved problem.

24 In the first part of the paper we address edge routing where an edge does
 25 not intersects the nodes it is not adjacent to. Our approach works for any node
 26 layout, as long id does not produce node overlaps. We build on the edge routing
 27 from [3] and improve it. There has been progress in visualization of large graphs
 28 recently. Tools appeared that can render a huge graph in seconds. However, the
 29 situation changes if we request that the node labels are rendered, and the edges
 30 overlap only the nodes they are adjacent to. Interacting with a large graph in an

31 Internet browser with the same ease as browsing an online map, inspecting the
32 high level structure and zooming in to the high level detail, is still an unsolved
33 problem. In this paper we describe novel approaches to several aspects of this
34 problem.

35 We propose a novel and efficient algorithm for edge routing, where each edge
36 can only intersect its source or target. The algorithm produces edge paths which
37 are visually appealing and even optimal in their homotopy class.

38 To facilitate graph visualization with WebGL, we propose a new simple and
39 fast tiling method. The method guarantees that in every view, except of the
40 views of the Shighest layer, the number of visible entities is not larger than a
41 predefined bound. The method can be used in other viewers that support tiling.

42 Our method provides a high level overview of the graph.

43 The edge routing algorithm mentioned above is reused at the tiling stage to
44 simplify the paths on the lower levels. In addition, we bundle edges per-tile as
45 an optimization heuristic. Our software runs calculations on the client desktop
46 or a phone, and renders the graph in the browser.

47 Related work

48 A popular graph drawing tool Graphviz [4] applies Scalable Force-Directed Place-
49 ment [5] for large graphs, with no support for tiling. Its edge routing for this case
50 builds the whole visibility graph. This can be very slow because the visibility
51 graph can have $O(n^2)$ edges, where n is the number of the nodes in the graph.
52 Interestingly, the funnel algorithm [6, 7], the last step of our approach, is used
53 in Graphviz for the edge routing in the Sugiyama layout. We are not aware of
54 any tool that integrates Graphviz and uses tiling as well.

55 yWorks [8] has method "Organic edge routing" that produces edge routes
56 around the nodes. We could find only a very general description of the method:
57 "The algorithm is based on a force directed layout paradigm. Nodes act as re-
58 pulsive forces on edges in order to guarantee a certain minimal distance between
59 nodes and edges. Edges tend to contract themselves. Using simulated annealing,
60 this finally leads to edge layouts that are calculated for each edge separately".
61 It seems the algorithm runs in $O(n + m) \log(n + m)$ time, where n is the number
62 of the nodes and m is the number of the edges.

63 ReGraph [9] uses WebGL as the viewing platform. It can render a large graph
64 using straight lines for the edges. The tool does not support tiling, but instead
65 the user interactively opens the node that is a cluster of nodes.

66 "graph-tool.skewed" [10] does not implement its own layout algorithms or
67 edge routing algorithms, but instead provides a nice wrapper around the algo-
68 rithms from other layout tools.

69 Circos [11] visualizes large graphs in a circular layout. It does not support
70 tiles.

71 Cosmograph [12] uses a GPU to calculate the layout of a graph and can
72 handle a graph with a million nodes. It renders edges as straight lines. It does
73 not support tiling.

74 The authors of [13] implemented GraphMaps, a tool for large graph visual-
 75 ization. The tool only ran in Windows. The edge were routed as polylines on a
 76 triangulation. The tool supported tiling, but the problem of the limiting number
 77 of visible entities was not solved.

78 In [14] an approach to visualize a huge graph. The method uses tiles and
 79 edge bundling following [15], which is applied at the last moment during the
 80 graph browsing. The latter calculation is done on the client side, the rest of the
 81 calculations run on several server machines.

82 Edge routing in XJS

93 The edge routing starts, as in [3], by building a spanner graph, an approximation
 94 of the full visibility graph, and then finding shortest paths on the spanner. The
 95 spanner, see Fig 2, is built on a variation of a Yao graph, which was introduced
 96 independently by Flinchbaugh and Jones [16] and Yao [17]. This kind of graph
 97 is defined by the set of cones with the apices at the vertices. The cones have the
 98 same angle, usually in the form of $\frac{2\pi}{n}$, where n is a natural number. The family
 99 of cones with the apex at a specific vertex partition the plane as illustrated in
 100 Fig. 1. For each cone at most one edge is created connecting the cone apex with
 101 a vertex inside of the cone, so the graph has $O(n)$ edges where n is the number
 102 of vertices.

103
 107 The approach of [3] applies local optimizations to shorten an edge path.
 108 Namely, it tries to shortcut one vertex at a time from the path, as illustrated in
 109 Fig 3. To smoothen a path, it fits Bezier segments into the polyline corners by
 110 using a binary search to find a larger fitting segment, see Fig 4. While analyzing
 111 performance of the edge routing in XJS, we noticed, that for a graph with more
 112 than 1k nodes these heuristics sometime create a performance bottleneck, in
 113 spite of using R-Trees[18].

114 In addition, when the naive shortcutting of polyline corners fails, the resulting
 115 path might remain not visually appealing, as shown in Fig. 3.

116 We replace these heuristics with a more precise and efficient optimization
 117 described below.

118 Path optimization

119 We finalize edge routes by the “funnel” algorithm [6, 7], routing a path inside a
 120 simple polygon, that is a polygon without holes.

121 An application of the ‘path in a simple polygon’ optimization to edge routing
 122 is not a new idea: the novelty of our work is in how we find the polygon and
 123 how we use it. The authors of Graphvis used the ‘funnel’ algorithm [19], but
 124 only for hierarchical layouts, where a simple polygon, \mathcal{P} , containing the path is
 125 available. They write: “If \mathcal{P} does not contain holes ... we can apply a standard
 126 “funnel” algorithm ... for finding Euclidean shortest paths in a simple polygon”.



Fig. 1. Yao graph




Fig. 2. Spanner graph is built using the idea of Yao graphs. The dashed curves are the original node boundaries. Each original curve is surrounded by a polygon with some offset to allow the polyline paths smoothing without intersecting the former.

The edge marked by the circles is created because the top vertex is inside of the cone and it is the closest among such vertices to the cone apex. The apex of the cone is the lower vertex of the edge.

XJS uses cone angle $\frac{\pi}{6}$, so the edges of the spanner can deviate from the optimal direction by this angle. Therefore, the shortest paths on the spanner have length that is at most the optimal shortest length multiplied by $\frac{1}{\cos(\frac{\pi}{6})} \simeq 1.155$.



Fig. 3. Unsuccessful shortcut

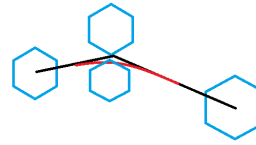





Fig. 4. Fitting a Bezier segment into a polyline corner




138 **Fig. 5.** Path \mathcal{L} with \mathcal{T} , a fragment.



139 **Fig. 6.** Polygon \mathcal{P} containing \mathcal{L} .



140 **Fig. 7.** New triangulation of \mathcal{P} .



141 **Fig. 8.** The optimized path together
142 with the sleeve diagonals.

127 In general case, for a non-layered layout, they build the visibility graph which is
128 very expensive for a large graph.

129 Here we find the polygon \mathcal{P} for any layout. We drop the requirement that
130 \mathcal{P} is simple. Indeed, to run the “funnel” algorithm one only needs a “sleeve”: a
131 sequence of triangles leading from the start to the end of the path, where each
132 triangle shares a side with its successor. Let us show how to build polygon \mathcal{P} ,
133 create a sleeve, and produce an optimized path.

134 We call obstacles, \mathcal{O} , the set of polygons covering the original nodes, see
135 Fig. 2. Before routing edges, we calculate a Constrained Delaunay Triangulation
136 [20] on \mathcal{O} . Let us call this triangulation \mathcal{T} .

137 For each edge of the graph we proceed with the following steps.

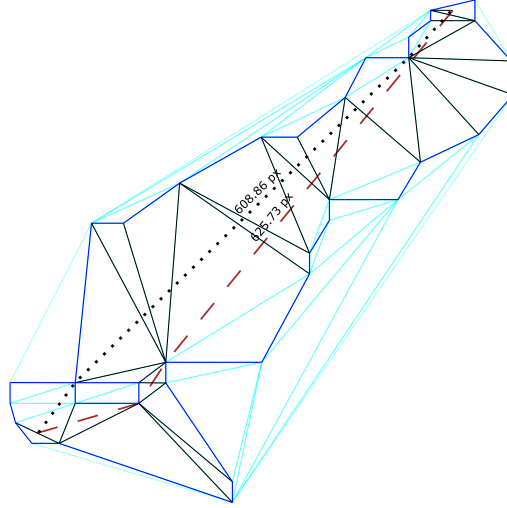
145 We route a path, called \mathcal{L} , on the spanner, as illustrated by Fig. 5. Let \mathcal{S} and
146 \mathcal{E} be the obstacles containing correspondingly \mathcal{L} ’s start and end point. To obtain
147 \mathcal{P} , let us consider \mathcal{U} , the set of all triangles $t \in \mathcal{T}$ such that either $t \subset \mathcal{S} \cup \mathcal{E}$, or t
148 intersects \mathcal{L} and is not inside of any obstacle in $\mathcal{O} \setminus \{\mathcal{S}, \mathcal{E}\}$. The union of \mathcal{U} gives
149 us \mathcal{P} . The boundary of \mathcal{P} comprizes all sides e of the triangles from \mathcal{U} such that
150 e belongs to exactly one triangle from \mathcal{U} , see Fig. 6.

151 To create the sleeve [6, 7], we need to have a triangulation of \mathcal{P} such that every
152 edge of the triangulation is either a boundary edge of \mathcal{P} , or a diagonal of \mathcal{P} .
153 Because \mathcal{U} might not have this property, as in Fig. 6, we create a new Constrained
154 Delaunay Triangulation of \mathcal{P} , where the set of constrained edges is the boundary
155 of \mathcal{P} , see Fig. 7.

156 We trace path \mathcal{L} through the new triangulation and obtain the sleeve. Finally,
157 we apply the funnel algorithm on the sleeve and obtain the path which is the
158 shortest in the homotopy class of \mathcal{L} , as illustrated in Fig. 8.

159 The discussion [21] of the algorithm helped us in the implementation.

160 Polygon \mathcal{P} is not necessarily simple, as shown in Fig. 9. In this example the
161 path that we calculate with the funnel algorithm is not the shortest path inside
162 of \mathcal{P} .



143 **Fig. 9.** \mathcal{P} is not simple. The dotted path is shorter than the dashed one that
 144 was found by the routing.

163 Performance and quality comparison

167 In Fig. 10 we compare the paths generated by the old and the new method. We
 168 can see that the paths produced by the new method have no kinks. We also
 169 know that these paths are the shortest in their 'channels'. Arguably, the new
 170 method produces better paths.

171 Our performance experiments are summarized in Table. 1. We see that the
 172 older approach outperforms the new one on the smaller graphs; those with the
 173 number of nodes under 2000. The new method is faster on the rest of the graphs.
 174 We still prefer to use the new method independently of the graph size since
 175 the slowdown is insignificant, but the quality of the paths is better. On the
 176 larger graphs the new method runs faster and produces better paths, so it is an
 177 obvious choice. To load a large graph, for example, `deezer_europe_edges` [22], we
 178 start Edge or Chrome with an option that increases the memory limit of their
 179 process: `- max_old_space_size=8192`.

191 1 Tiling

192 The algorithm works in three phases. The first phase builds the levels starting
 193 from the lowest level and proceeding to higher and more detailed levels, with
 194 smaller tiles, until no more tile subdivision is required. The second phase filters
 195 out the entities from the layers to satisfy the capacity quota. Finally, the third
 196 phase simplifies the edge routes to utilize the space freed by the filtered out
 197 entities.




Fig. 10. The difference in the paths between the old, on the left, and the new, on the right, paths. The arrows on the left fragment point to the kinks that were removed by the new method.

graph	nodes	edges	old method's time	new time
social network [23]	407	2639	1.0	1.4
b103 [24]	944	2438	1.6	2.0
b100 [25]	1463	5806	5.6	5.785
composers [26]	3405	13832	510.5	20.3
p2p-Gnutella04 [27]	10876	39994	375.4	304.2
facebook_combined [28]	4039	88234	132.2	123.7
lastfm_asia_edges [22]	7626	27807	43.3	54.7
deezer_europe_edges [22]	28283	92753	1596.9	1402.6
ca-HepPh [29]	12008	237010	521.2	495.0

Table 1. Performance comparison with time in seconds.

198 A tile, in our settings, is a pair $(rect, tiledata)$, where $rect$ is the rectangle of
 199 the tile and $tiledata$ is a set of *tile elements* visible in $rect$. A *tile element* could
 200 be a node, an edge label, an edge arrowhead, or an *edge clip*. An edge clip is a
 201 pair (e, p) , where e is an edge and p is a continuous piece of the edge curve c_e .
 202 Sometimes we need several edge clips to trace an edge through a tile.

203 The initial tile, the only tile on level 0, is represented by pair $(0, 0)$. For
 204 $z = 1$, there are four tiles: $(0, 0)$, $(0, 1)$, $(1, 0)$, and $(1, 1)$. Each tile (i, j) can be
 205 subdivided into four subtiles for level $z + 1$: $(2i, 2j)$, $(2i, 2j + 1)$, $(2i + 1, 2j)$, and
 206 $(2i + 1, 2j + 1)$.

207 Each z -level is represented by a map L_z , so $L_z(i, j)$ gives us a specific tile.
 208 Empty tiles correspond to undefined $L_z(i, j)$.

209 We use edge clips to represent the edge intersections with the tiles and provide
 210 the renderer with the minimal geometry that is sufficient to render a tile. To
 211 achieve this we require property \mathcal{F} :

- 212 a) For each tile t , for each curve clip $(e, p) \in t.tiledata$, we have: $p \subset t.rect$
 213 and p might cross the boundary of the $t.rect$ only at endpoints of p .
- 214 b) For each edge e we have : the union of all p for all $(e, p) \in t.tiledata$ is
 215 equal to $c_e \cap t.rect$.

216 First phase of tiling

217 The first phase starts with $L_0 = \{(0, 0) \rightarrow (rect, tiledata)\}$: and $tiledata$ com-
 218 prising curve clips (e, c_e) , for all edges e of the graph, all graph nodes, all edge
 219 labels, and all edge arrowheads. We ensure property \mathcal{F} by setting $rect$ to a
 220 padded bounding box of the graph, so each edge curve does not intersect the
 221 boundary of $rect$.


222 Let us assume that L_z is already constructed and \mathcal{F} holds for its tiles. To
 223 build level L_{z+1} we divide each tile $t = L_z(i, j)$ into four subtiles of equal size.
 224 For each node, arrowhead, or edge label of $t.tiledata$, if the bounding box of the
 225 element intersects the subtile's rectangle then we add the element to the subtile
 226 $tiledata$.

227 The edge clip treatment is more involved. Let (e, p) be a curve clip belonging
 228 to tile t . We find all intersections of curve p with the horizontal midline and the
 229 vertical midline of $t.rect$. Each intersection can be represented as $p[t_j]$. We sort
 230 sequence $u = [start, \dots, t_j, \dots, end]$, where $[start, end]$ is the parameter domain
 231 of p , in ascending order, and remove the duplicates.

232 Next we create curve clips $(e, l_k) = (e, trim(p, u_k, u_{k+1}))$, as shown in Fig 11.
 233 We assign each curve clip (e, l_k) to the subtile with the rectangle containing the
 234 bounding box of l_k .

238 Because, by the induction assumption property \mathcal{F} is true on L_z , and by
 239 construction, each new curve clip can cross the boundary of the subtile only at
 240 the clip endpoints. We also cover all the intersections of p with the subtiles with
 241 the new edge clips, so the property \mathcal{F} holds for L_{z+1} .

242 Two parameters control the algorithm: tile capacity, \mathcal{C} , and the minimal size
 243 of a tile: $(\mathcal{W}, \mathcal{H})$. If for each (i, j) the number of elements in $L_z(i, j).tiledata$ is



235 **Fig. 11.** Intersect curve $[start, end]$ with the midlines. Sort the intersections pa-
 236 rameters together with start, and end into array $u = [start, a, b, c, end]$. Split the
 237 curve to sub-curves $[start, a]$, $[a, b]$, $[b, c]$, $[c, end]$.

244 not greater than \mathcal{C} , or, if $w \leq \mathcal{W}$ and $h \leq \mathcal{H}$, where w (h) is the current tile
 245 width (correspondingly, height), then the second phase starts.

246 In our setting $\mathcal{C} = 500$, and $(\mathcal{W}, \mathcal{H}) = 3(w, h)$, where w is the average width
 247 and h is the average height of the nodes of the graph.

248 **Edge bundling** In our settings each edge clip is uniquely defined, module
 249 direction, by its start and end point. We can use this property to bundle the
 250 edges. In each tile we keep a map from unordered pairs of points to the set of
 251 edge clips that have these points as start and end points. Each such pair defines
 252 an edge bundle. For all edge clips in a bundle we create only one curve segment,
 253 avoiding the expensive trimming. We also count a bundle as one element in the
 254 tile, as in most of the cases the drawing attributes of the edges in the bundle are
 255 the same.

256 In our experiments, the number of edge bundles is about 50% of the number
 257 of edge clips, so the edge bundling is a significant optimization.

258 Second phase of tiling

259 The second phase of tiling filters out the entities from the lower layers. We do not
 260 change the highest, the most detailed layer. We sort the nodes of the graph into
 261 array N by PageRank [30]. For each layer L , except of the highest, we proceed
 262 as follows.

263 Here `removeEntities(L)` empties all the tiles of layer L , but returns map r allow-
 264 ing to restore the tiles. Function `addNodeToLayer(n)` returns false and does not
 265 change L when one of the tiles intersecting n already has more elements than
 266 \mathcal{C} . Otherwise, the function adds n to all tiles intersected by n . It also adds the
 267 tile elements for self edges of n , and the tile elements for the edges connecting n

```

1: procedure FILTER( $L$ )
2:    $r \leftarrow \text{removeEntities}(L)$ 
3:   for all  $n$  in  $N$  do
4:     if ! $\text{addNodeToLayer}(n, r, N)$  then break
5:     end if
6:   end for
7: end procedure

```

with the nodes appearing in N before n , i.e. the nodes with the rank not lesser than the rank of n .

This procedure guarantees that each tile of L has no more than \mathcal{C} nodes, but a tile can have more than \mathcal{C} elements in general.

Third phase of tiling

In the third phase we use a fact that some nodes are not present on the layer. For all layers, except of the highest, we reroute the edges but only around the nodes that are present in the layer. We do not calculate edge routes from scratch, but use the existing routes and only apply the "funnel" heuristic in larger channels. This gives us simpler edge routes but still has a visual stability during the layer change while browsing.

2 Future work

- Find a tiling method that guarantees that each tile has no more than \mathcal{C} elements. One approach could be to use a more aggressive edge bundling to reduce the number of edge clips in the tiles.

References

1. Y. Hu and L. Shi, "Visualizing large graphs," *Wiley Interdisciplinary Reviews: Computational Statistics*, vol. 7, no. 2, pp. 115–136, 2015.
2. U. Brandes and C. Pich, "Eigensolver methods for progressive multidimensional scaling of large data," in *Graph Drawing: 14th International Symposium, GD 2006, Karlsruhe, Germany, September 18-20, 2006. Revised Papers 14*, pp. 42–53, Springer, 2007.
3. T. Dwyer and L. Nachmanson, "Fast edge-routing for large graphs," in *Graph Drawing: 17th International Symposium, GD 2009, Chicago, IL, USA, September 22-25, 2009. Revised Papers 17*, pp. 147–158, Springer, 2010.
4. "Graphviz." <http://www.graphviz.org/>.
5. "sfdp." <https://graphviz.org/docs/layouts/sfdp/>.
6. B. Chazelle, "A theorem on polygon cutting with applications," in *23rd Annual Symposium on Foundations of Computer Science (sfcs 1982)*, pp. 339–349, IEEE, 1982.

- 298 7. J. Hershberger and J. Snoeyink, "Computing minimum length paths of a given
299 homotopy class," *Computational geometry*, vol. 4, no. 2, pp. 63–97, 1994.
- 300 8. "yworks." <https://yworks.com/products/yed>.
- 301 9. "Regraph." <https://cambridge-intelligence.com/regraph/>.
- 302 10. "Skewed." <https://graph-tool.skewed.de>.
- 303 11. "Circos." <http://circos.ca/>.
- 304 12. "Cosmograph." <https://cosmograph.app>.
- 305 13. L. Nachmanson, R. Prutkin, B. Lee, N. H. Riche, A. E. Holroyd, and X. Chen,
306 "Graphmaps: Browsing large graphs as interactive maps," in *Graph Drawing and*
307 *Network Visualization: 23rd International Symposium, GD 2015, Los Angeles, CA,*
308 *USA, September 24-26, 2015, Revised Selected Papers 23*, pp. 3–15, Springer, 2015.
- 309 14. A. Perrot and D. Auber, "Cornac: Tackling huge graph visualization with big data
310 infrastructure," *IEEE Transactions on Big Data*, vol. 6, no. 1, pp. 80–92, 2018.
- 311 15. C. Hurter, O. Ersoy, and A. Telea, "Graph bundling by kernel density estimation,"
312 in *Computer graphics forum*, vol. 31, pp. 865–874, Wiley Online Library, 2012.
- 313 16. B. Flinchbaugh and L. Jones, "Strong connectivity in directional nearest-neighbor
314 graphs," *SIAM Journal on Algebraic Discrete Methods*, vol. 2, no. 4, pp. 461–463,
315 1981.
- 316 17. A. C.-C. Yao, "On constructing minimum spanning trees in k-dimensional spaces
317 and related problems," *SIAM Journal on Computing*, vol. 11, no. 4, pp. 721–736,
318 1982.
- 319 18. A. Guttman, "R-trees: A dynamic index structure for spatial searching," in *Pro-*
320 *ceedings of the 1984 ACM SIGMOD international conference on Management of*
321 *data*, pp. 47–57, 1984.
- 322 19. D. P. Dobkin, E. R. Gansner, E. Koutsofios, and S. C. North, "Implementing a
323 general-purpose edge router," in *Graph Drawing: 5th International Symposium,*
324 *GD'97 Rome, Italy, September 18–20, 1997 Proceedings 5*, pp. 262–271, Springer,
325 1997.
- 326 20. B. Delaunay, "Sur la sphere vide, bull. acad. science ussr vii: Class," *Sci. Mat. Nat*,
327 pp. 793–800, 1934.
- 328 21. "Funnel algorithm." <https://page.mi.fu-berlin.de/mulzer/notes/alggeo/polySP.pdf>.
- 329 22. B. Rozemberczki and R. Sarkar, "Characteristic Functions on Graphs: Birds of a
330 Feather, from Statistical Descriptors to Parametric Models," in *Proceedings of the*
331 *29th ACM International Conference on Information and Knowledge Management*
332 *(CIKM '20)*, p. 1325–1334, ACM, 2020.
- 333 23. A. Beveridge and M. Chemers, "The game of game of thrones: Networked con-
334 cordances and fractal dramaturgy," in *Reading Contemporary Serial Television*
335 *Universes*, pp. 201–225, Routledge, 2018.
- 336 24. "b103." [https://github.com/microsoft/automatic-graph-](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot)
337 [layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b103.dot).
- 338 25. "b100." [https://github.com/microsoft/automatic-graph-](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot)
339 [layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot](https://github.com/microsoft/automatic-graph-layout/blob/master/GraphLayout/Test/MSAGLTests/Resources/DotFiles/LevFiles/b100.dot).
- 340 26. "Skewed." <http://mozart.diei.unipg.it/gdcontest/contest2011/composers.xml>.
- 341 27. "p2p-gnutella04." <https://snap.stanford.edu/data/p2p-Gnutella04.html>.
- 342 28. "facebookcombined." https://snap.stanford.edu/data/facebook_combined.txt.gz.
- 343 29. J. Leskovec, J. Kleinberg, and C. Faloutsos, "Graph evolution: Densification
344 and shrinking diameters," *ACM transactions on Knowledge Discovery from Data*
345 *(TKDD)*, vol. 1, no. 1, pp. 2–es, 2007.
- 346 30. L. Page, S. Brin, R. Motwani, and T. Winograd, "The pagerank citation ranking:
347 Bringing order to the web," *Stanford InfoLab*, vol. 249, no. 373, pp. 1–17, 1999.